

Wrangling Rogues: A Case Study on Managing Experimental Post-Moore Architectures

Will Powell, Jason Riedy, Jeffrey Young, Tom Conte

1 August 2019

Center for Research into Novel Computing Hierarchies at Georgia Tech

Outline

What is the CRNCH Rogues Gallery?

Current Rogues

- Emu Chick

- 3D Stacked Memories and FPGAs

- Neuromorphic / Analog Hardware (FPAA)

Management lessons learned

- Helpful points

- Painful points

Apps: Massive+-scale data analysis

Cyber-security

Identify anomalies, malicious actors

Health care

Find outbreaks, population epidemiology, similar patient association

Social networks

Advertising, searching, grouping

Intelligence

Decisions at scale, regulating markets, smart & sustainable cities

Systems biology

Understanding interactions, drug design

Power grid / Smart cities

Disruptions, conservation, prediction

Irregular data access. Changing data.

The New York Times
Thursday, September 4, 2008

Wrangling Rogues — 1 Aug 2019

By MATTHEW L. WALD
Published: November 12, 2003

E-MAIL

Massive Social Network Analysis:
Mining Twitter for Social Good

10 @backchannel
11 @TIME
12 @CDCemergency
13 @CDC_eHealth
14 @perezhillton
15 @billmaher

@JTsDJs
@ATLien
@MarshallRams
@Kanye

3/22

High-Performance Data Analysis (HPDA)

Novel applications:

- Data at scale and speed needs new ideas for computing analysis.
- “Big data” platforms fare poorly v. a single thread plus large SSD even for static data sets. (McSherry, Isard, Murray. “Scalability! But at what COST?” HotOS XV, 2015.)
- Many high-level codes are written and re-written to answer one question: need flexibility.
- *Some* primitives may be tuned and re-used.


Why do we need rogues?

Rogue: Someone who goes their own way, who breaks away from the crowd.



- Current architectures are hitting limits on manufacturing, heat dissipation, memory latency...
- What happens when novel prototypes hit reality?
- Designers need feedback, a software ecosystem, and **trained students.**

What is the Rogues Gallery?



Hardware! “I’ll tell you later.”

Introducing the CRNCH Rogues Gallery

CRNCH Rogues Gallery

A physical & virtual space for hosting novel computing architectures, systems, and accelerators since fall 2017.

Host / manage remote access for novel architectures to

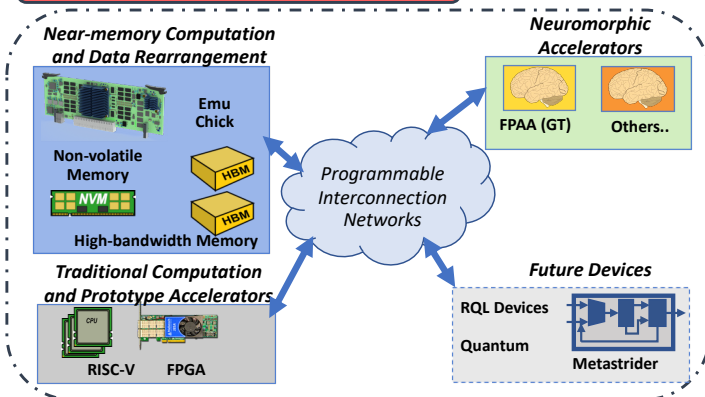

- kick-start software ecosystems (*e.g.* Kokkos),
- leverage real applications to train students, and
- provide rapid feedback to architects.

Amortize effort and cost of trying novel architectures.

Break the “but it’s too much work” barrier.

<http://crnch.gatech.edu/rogues-gallery>

Rogues Gallery summary



Current Rogues

Current Rogues


Emu Chick

3D Stacked Memories and FPGAs


Neuromorphic / Analog Hardware (FPAA)

Emu Technology's Chick

- “Migratory Memory Side Processing” to exploit *weak locality*.
- Data for graph edge attributes, documents / medical records, etc. reside nearby even if accessed irregularly.
- Moving threads to data on reads: all reads are local, so lower latency.




Emu Chick: Ongoing Results



Platform:

- Application direction.
- Demo-able code.
- Debugging.
- Connections.



People:



- GT → Emu: One employee, one intern.
- Emu → GT: One graduate student.

3D Stacked Memory and FPGAs

- FPGA + HMC / DRAM: Enable experiments with “near-memory” and memory-centric processing.
- FPGA platforms prototype non-traditional accelerators like Automata, sparse data engines, *etc.*
- Current work is supported in part by Micron hardware donation.



FPGA & memory results




Hadidi, Asgari, Young, Mudassar, Garg, Krishna, Kim. "Performance Implications of NoCs on 3D-Stacked Memories: Insights from the Hybrid Memory Cube (HMC)," ISPASS 2018



- Characterizations with FPGA and Hybrid Memory Cube show latency/bandwidth tradeoff.
- Other FPGA work is focused on compilers, HPC prototyping, and sparse algorithms for Intel and Xilinx FPGAs.

Neuromorphic systems

- Field-Programmable Analog Array (FPAA) System-On Chip, designed in the lab of Dr. Jennifer Hasler.
- Analog + digital to achieve unprecedented power and size reductions.
- Potential on-chip/package accelerator.
- Adding other neuromorphic systems



FPAA results



Embedded learning & classification: 20-30 μ W
on full, 1s Nzero database (GOMAC 2016)

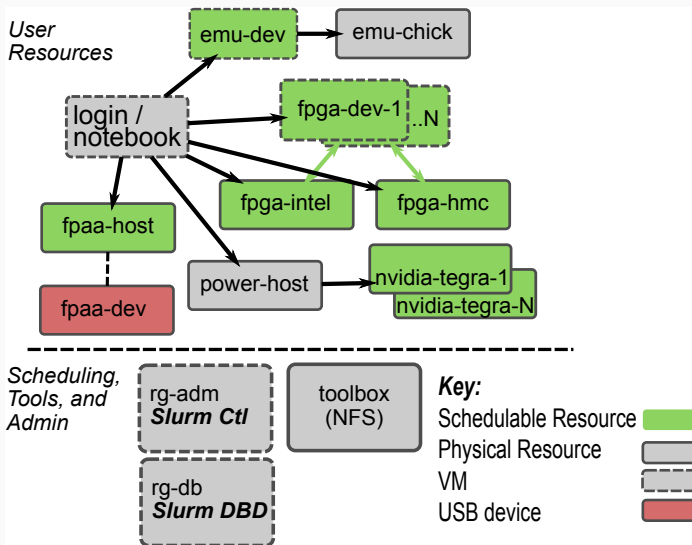
Management lessons learned

Management lessons learned

Helpful points

Painful points

Rogues Gallery structure



Management lessons learned

- Invest in rogues, but realize some technology may be short-lived.
 - Minimize custom management effort.
- Physical hardware resources not dedicated to rogues should be kept to a minimum.
 - Don't spend \$ on non-rogues.
- Collaboration and commiseration is key.
 - Rogues need a community to succeed.
- Licensing and appropriate identity management are tough but necessary challenges.
 - Use network isolation when needed.

Helpful points

- Network isolation provides security.
 - Well, enough given limited usefulness.
- Singularity is great for build environments.
 - HW start-ups cannot afford supporting every OS/arch.
 - IT cannot afford supporting every OS/arch.
 - Companies must be friendly...
- Inspired undergrads are wonderful!
 - Modernizing tools (FPAA)
 - Building out demonstrations
 - <http://www.vip.gatech.edu/teams/rogues-gallery>

Painful points

- SLURM aspects:
 - Managing slurmd.conf.
 - Building on all the OS/arch combos.
- Few light-weight management options.
 - salt-ssh, ansible on *some*
- Hardware access for rebooting, reseating.
- Many programming interfaces, few people
 - Kokkos, TENNLab, more...
- Still need to tackle “sensitive” data, including some FPGA IP
- Reproducible / replicable / audit-able results

Rogues Gallery: Active and Growing

- Integrating FPAAs and toolchain
- Tight development loop with Emu
- Active research projects and publications
- Community building via tutorials & talks
- New approaches to benchmarking, quantum software stacks, neuromorphic toolchains, ...

CRNCH Rogues Gallery connects researchers *and students* with novel architectures and architects with upcoming applications.

Let us host / manage your neat stuff!

<http://crnch.gatech.edu/rogues-gallery>

Acknowledgments

Fantastic students and colleagues:

- Srinivas Eswar (GT CSE)
- Dr. Eric Hein (GT ECE \Rightarrow Emu)
- Patrick Lavin (GT CSE)
- Dr. Jiajia Li (GT CSE \Rightarrow PNNL)
- Abdurrahman Yaşar (GT CSE)
- Chunxing Yin (GT CSE)
- Dr. Jeffrey S. Young (GT CS)
- Dr. Tom Conte (GT CS/ECE)
- Dr. Vivek Sarkar (GT CS)
- Dr. Ümit Çatalürek (GT CSE)
- Dr. Bora Uçar (ENS Lyon CNRS)
- Dr. Rich Vuduc (GT CSE)

Code (ideally will have links from `crnch.gatech.edu`):

- <https://gitlab.com/crnch-rg>
- <https://github.com/ehein6/emu-microbench>

Other testbeds:

- ORNL: ExCL
- Argonne
- Berkeley: AQCT
- PNNL: CENATE
- Sandia HAAPS
- (others?)

External Image Credits

- “What’s that watermelon doing there?”: copyright MGM, used for identification
- Oscar Wilde: public domain, obtained from Wikipedia
- Edna St. Vincent Millay: public domain, obtained from Wikipedia
- Dread Pirate Roberts: copyright 20th Century Fox, used for identification
- Mary Jackson, Katherine Goble Johnson, Dorothy Vaughan (Hidden Figures): copyright 20th Century Fox, used for identification
- Malcolm Reynolds: copyright Universal Pictures, used for identification
- Rogue One: copyright Walt Disney Studios Motion Pictures, used for identification
- The Story of Karrawingi, the Emu (cover): copyright estate of Leslie Rees, used for identification
- Big Hero 6: copyright Walt Disney Studios Motion Pictures, used for identification