

# Jupyter-Based Self-Service Training in OOD for Novel Compute Architectures

A. Jezghani

Georgia Institute of Technology, Atlanta, GA



# Who We Are



Aaron Jezghani  
PACE



Kevin Manalo  
PACE (former)



William Powell  
CSE



Jeffrey Valdez  
IC



Jeffrey Young  
CS

The background of this slide is a grayscale aerial photograph of a dense urban area with numerous buildings and infrastructure. A red rectangular box is overlaid on the lower-left portion of the image, containing text about the CoC - CCB. In the top right corner of the slide, there is a white box with a teal border containing text about PACE - Coda.

### CoC - CCB

- Ubuntu and RHEL
- AMD, Intel, Power9, Arm
- Ethernet, IB, OPA
- GPU, FPGA, FPAA, DPU, etc.

### PACE - Coda

- RHEL
- Intel
- IB
- GPU

Google 100% Imagery date: 1/16/20 Landsat / Copernicus

80 m Camera: 716 m 33°46'24"N 84°23'38"W 289 m



Introduction



Rogues  
Gallery



Challenges



Deploying on  
A64FX



Jupyter  
Workflow

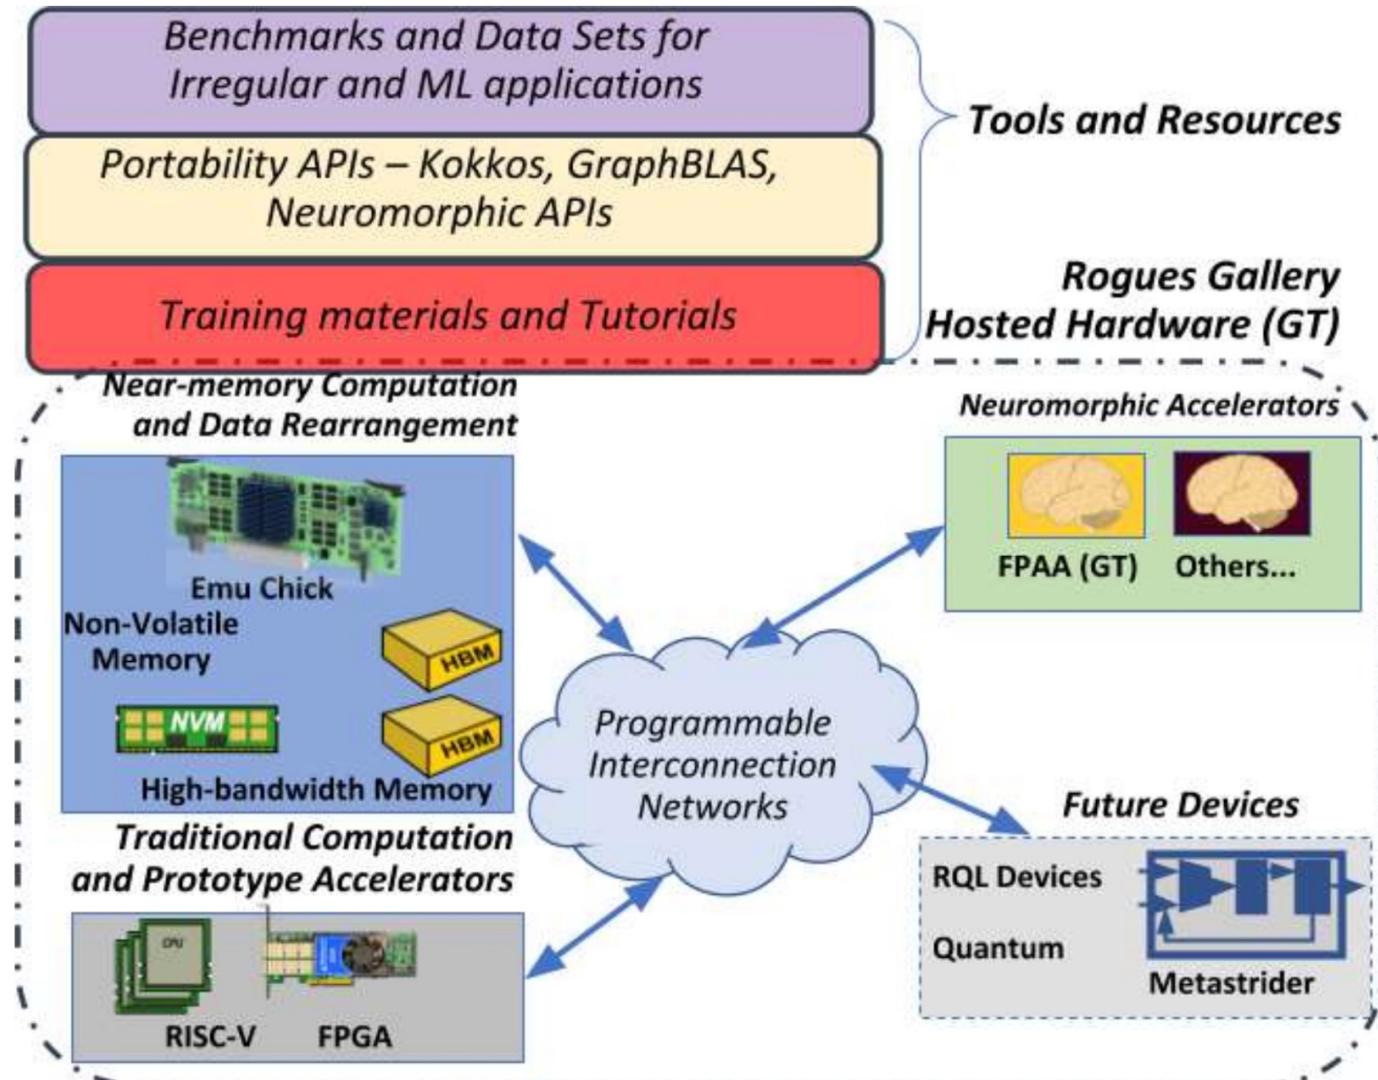


Summary

# The Rogues Gallery

**NSF MRI award #1828187:** “MRI: Acquisition of an HPC System for Data-Driven Discovery in Computational Astrophysics, Biology, Chemistry, and Materials Science.”

**NSF award #2016701:** “CCRI: Medium: Rogues Gallery: A Community Research Infrastructure for Post-Moore Computing”



# The Hurdles of New Technologies and Users

- As exciting as algorithm exploration and software porting might be for veteran users...
  - The “tried-and-true” tools often take several cycles to adapt
    - This happens with updates to traditional architectures too
    - Vendors contribute heavily to upstream, but it can still take time to trickle down
  - Vendor-provided tools are great to use!
    - ...but even as a port of something familiar, there’s still a learning curve
    - And these tools may be proprietary!
- Meanwhile, new users are constantly inbound
  - Remote command line access can be intimidating (and feel prohibitive)
  - Even for experienced users, interactive server workflows can be problematic



# These Aren't New Issues! But...

- Knowledge transfer can be a difficult process
  - Documentation relies on the "Goldilocks principle"...on a user-by-user basis.
  - Trainings are great but might not be convenient for the users.
  - Command-line self-service tutorials may not include the instructive element available in other forms.
- Many solutions are less than elegant
  - Performance problems (X11 forwarding)
  - Scalability issues (FastX licensing, X2Go resource needs)
  - Less-than intuitive interfaces (SSH tunneling)
  - Client application requirements (e.g. VNC)
  - Local security rules can be an impediment (Port blocking)



# Configuring Compute Nodes for OOD

- Not surprisingly, it worked out of the box
  - (since our OOD VM is x86)
- Followed OSC templates that leverage xfce + TurboVNC
  - Aarch64 xfce is available from EPEL, but libturbojpeg needs to be manually built with the appropriate flag
- Everything else pretty much out of the box
  - Given the nature of the user community, the only modular software we've provided at the moment are compilers, math libraries, and MPI



# How We Leverage Jupyter

- Apps directory set up by architecture
  - E.g. all Octavius-related apps prefixed with a64fx\_\*
  - Access to apps controlled via POSIX group, e.g. rg-arm-users
- We provide a dropdown that auto-populates available templates
  - E.g., containerized versions of TF and PyTorch
- Tutorials are merely a port of the standard Jupyter notebook
  - Resource request fixed to match the needs of the tutorial
    - 8 cores on fixed queue and fixed time
  - Job script copies a template notebook to session directory
  - Key with tutorial design is to keep code execution native to mimic typical workflow while leveraging the ipython capabilities for enhancement

Introduction

Rogues  
Gallery

Challenges

Deploying on  
A64FX

Jupyter  
Tutorials

Summary

# Launching Containerized Notebooks

## form.yml.erb

```
<%-  
require "csv"  
output= File.read('/net/netscratch/ondemand-jupyter/templates.txt')  
templates = CSV.parse(output)  
-%>  
# Batch Connect app configuration file  
#  
# @note Used to define the submitted cluster, title, description, and  
# hard-coded/user-defined attributes that make up this Batch Connect app.  
---  
title: "Jupyter"  
cluster: "crnch"  
attributes:  
  modules: ""  
  extra_jupyter_args: ""  
bc_num_cores:  
  help: 'Number of CPUs/Cores.'  
  id: bc_num_cores  
  label: CPUs/Cores Per Node  
  max: 48  
  min: 1  
  step: 1  
  value: 1  
  widget: number_field  
  cacheable: false
```

```
custom_template:  
  label: "Starting Interface/Template"  
  widget: select  
  options:  
    - [ "None - use Jupyter to select file to work on", "" ]  
    - [ "This isn't real, because why not", "" ]  
<%- templates.each do |a| -%>  
  - [ "<%= a[0] %>", "<%= a[1] %>" ]  
<%- end -%>  
custom_queue:  
  widget: select  
  id: bc_queue  
  label: "Octavius Partition"  
  cacheable: false  
  options:  
    - [ "Debug: up to 4 nodes and 1 hour", "rg-arm-debug" ]  
    - [ "Short: up to 16 nodes and 2 hours", "rg-arm-short" ]  
    - [ "Long: up to 16 nodes and 6 hours", "rg-arm-long" ]  
form:  
  - extra_jupyter_args  
  - modules  
  - custom_queue  
  - bc_num_hours  
  - bc_num_slots  
  - bc_num_cores  
  #- bc_email_on_started  
  - custom_template
```



# Launching Containerized Notebooks

## submit.sh.erb

```
# Launch the Jupyter Notebook Server
<%- if context.custom_template == "TensorFlow.ipynb" -%>
  set -x
  export ONDEMAND_FILE="${odir}/<%=context.custom_template%>"
  export TEMPLATE_FILE="${sdir}/<%= context.custom_template %>"
  # Launch the Jupyter Notebook Server
  cp $ONDEMAND_FILE $TEMPLATE_FILE
  TEMPLATE_FILE=${TEMPLATE_FILE#*ondemand/}
  singularity run "${odir}"/tensorflow-arm.sif jupyter notebook --debug --
config="${CONFIG_FILE}" <%= context.extra_jupyter_args %> \
  --NotebookApp.default_url="/tree/ondemand/$TEMPLATE_FILE"
<%- elsif context.custom_template == "PyTorch.ipynb" -%>
  set -x
  export ONDEMAND_FILE="${odir}/<%=context.custom_template%>"
  export TEMPLATE_FILE="${sdir}/<%= context.custom_template %>"
  # Launch the Jupyter Notebook Server
  cp $ONDEMAND_FILE $TEMPLATE_FILE
  TEMPLATE_FILE=${TEMPLATE_FILE#*ondemand/}
  singularity run "${odir}"/pytorch-arm-neoverse-n1_r21.10-torch-1.9.0-openblas-jupyter.sif
jupyter notebook --debug --config="${CONFIG_FILE}" <%= context.extra_jupyter_args %> \
  --NotebookApp.default_url="/tree/ondemand/$TEMPLATE_FILE"
]
<%- else -%>
  set -x
  jupyter notebook --config="${CONFIG_FILE}" <%= context.extra_jupyter_args %>
<%- end -%>
```

# Arm Compiler Tutorial

## script.sh.erb

```
#!/usr/bin/env bash

# Benchmark info
echo "TIMING - Starting main script at: $(date)"

# Get Running Script Directory
get_script_dir () {
    SOURCE="${BASH_SOURCE[0]}"
    while [ -h "$SOURCE" ]; do
        DIR="$( cd -P "$( dirname "$SOURCE" )" && pwd )"
        SOURCE="$(
            readlink "$SOURCE"
            [[ $SOURCE != /* ]] && SOURCE="$DIR/$SOURCE"
        )"
    done
    $( cd -P "$( dirname "$SOURCE" )" )
    pwd
}
sdir="$(get_script_dir)"
export SDIR=$sdir
odir="/net/projects/notebook/ondemand-jupyter/"

# Set working directory to home directory
cd "${HOME}"

#
# Start Jupyter Notebook Server
#
<%- unless context.modules.blank? -%>
# Purge the module environment to avoid conflicts
module purge

# Load the require modules
module load <%= context.modules %>

# List loaded modules
module list
<%- end -%>

# Benchmark info
echo "TIMING - Starting jupyter at: $(date)"

# Launch the Jupyter Notebook Server
set -x
export ONDEMAND_FILE="${odir}/arm_compilers_evaluation.ipynb"
export TEMPLATE_FILE="${sdir}/arm_compilers_evaluation.ipynb"
cp $ONDEMAND_FILE $TEMPLATE_FILE
TEMPLATE_FILE=${TEMPLATE_FILE##*ondemand/}
jupyter notebook --config="${CONFIG_FILE}" <%= context.extra_jupyter_args %> \
--NotebookApp.default_url="/tree/ondemand/$TEMPLATE_FILE"
```



# Arm Compiler Tutorial

jupyter arm\_compilers\_evaluation (autosaved)

In [5]:

```
%bash --err naive
ml arm21
cd arm-sve-tools/01_Compiler/01_Naive
for compiler in {arm,cray-gnu}
do
    make clean >/dev/null
    make run COMPILER=${compiler} 2>&1
done | tee /dev/stderr 2>(sed -n 's/&*multiply took: ([0-9.]*).*&\1&p' >&2)

armclang --version
Arm C/C++/Fortran Compiler version 21.1 (build number 1069) (based on LLVM 12.0.1)
Target: aarch64-unknown-linux-gnu
Thread model: posix
InstalledDir: /net/projects/tools/aarch64/rhel-8/arm-allinea/21.1/arm-linux-compiler-21.1_Generic-AArch64_RHEL-8_aarch64-linux/bin
armclang++ -Rpass=\(loop-vectorize\) -Rpass-missed=\(loop-vectorize\) -o mm_arm_def.exe mm.cpp
-----
armclang --version
Arm C/C++/Fortran Compiler version 21.1 (build number 1069) (based on LLVM 12.0.1)
Target: aarch64-unknown-linux-gnu
Thread model: posix
InstalledDir: /net/projects/tools/aarch64/rhel-8/arm-allinea/21.1/arm-linux-compiler-21.1_Generic-AArch64_RHEL-8_aarch64-linux/bin
armclang++ -Rpass=\(loop-vectorize\) -Rpass-missed=\(loop-vectorize\) -Ofast -mcpu=native -o mm_arm_opt.exe mm.cpp
mm.cpp:42:13: remark: loop not vectorized [-Rpass-missed=sve-loop-vectorize]
    for (k = 0; k < m; ++k){
```

Read the results into a dataframe

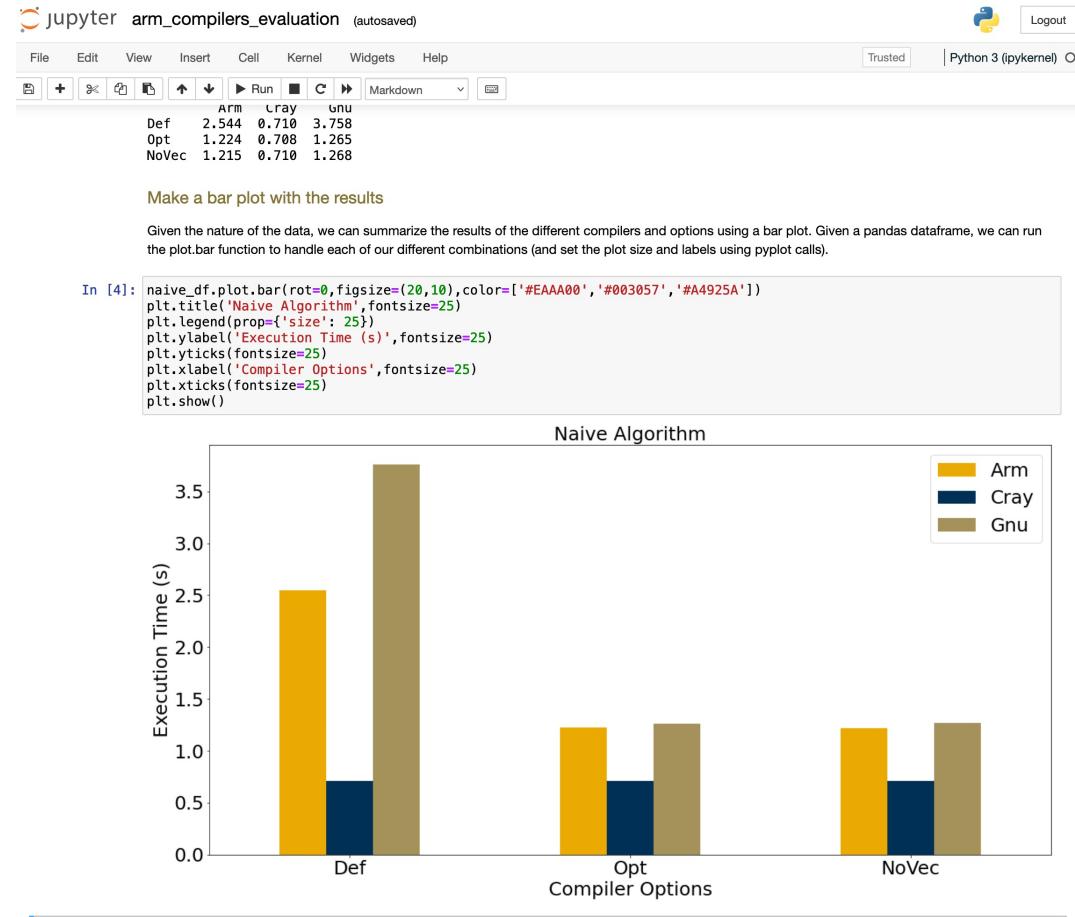
The values we stored in 'naive' are stored as a new-line delimited string, so we will use numpy.fromstring to read them into an array, reshape into a 3x3 2D array, transpose, and then create a dataframe with appropriate index and column labels. The restructuring of the data is so that we can take advantage of the "simpler" plotting calls in pandas versus matplotlib.

In [6]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

naive_df = pd.DataFrame(np.transpose(np.fromstring(naive,dtype=float,sep='\n')).reshape(3,3)),
                        index=['Def','Opt','NoVec'],columns=['Arm','Cray','Gnu'])
print(naive_df)
```

	Arm	Cray	Gnu
Def	2.582	0.710	3.758
Opt	1.224	0.708	1.268
NoVec	1.225	0.710	1.269



# Conclusion

- Onboarding new users can be a challenge
- OOD is a powerful platform to lower the barrier to entry
  - Provides full cluster access while improving GUI/interactive experiences
- Can leverage applications and example code to produce self-service tutorials
- Similarly, can leverage the dynamic content capabilities to abstract complex workflows (such as containerized solutions)
- Expanding Rogues Gallery OOD capabilities and looking to grow the community
  - Encourage innovative contributions to A64FX using Octavius

**Bored today/tomorrow? Check out the [CRNCH Summit 2022 videos](#)**

