

Facilitating Instruction on Exotic Architectures

Aaron Jezghani, PhD
Georgia Institute of Technology
March 19, 2024

NSF ACCESS All Champions Meeting

The Rogues Gallery



The Rogues Gallery is a **disaggregated testbed** consisting of hardware, VMs, and services hosted across 3 GT buildings.

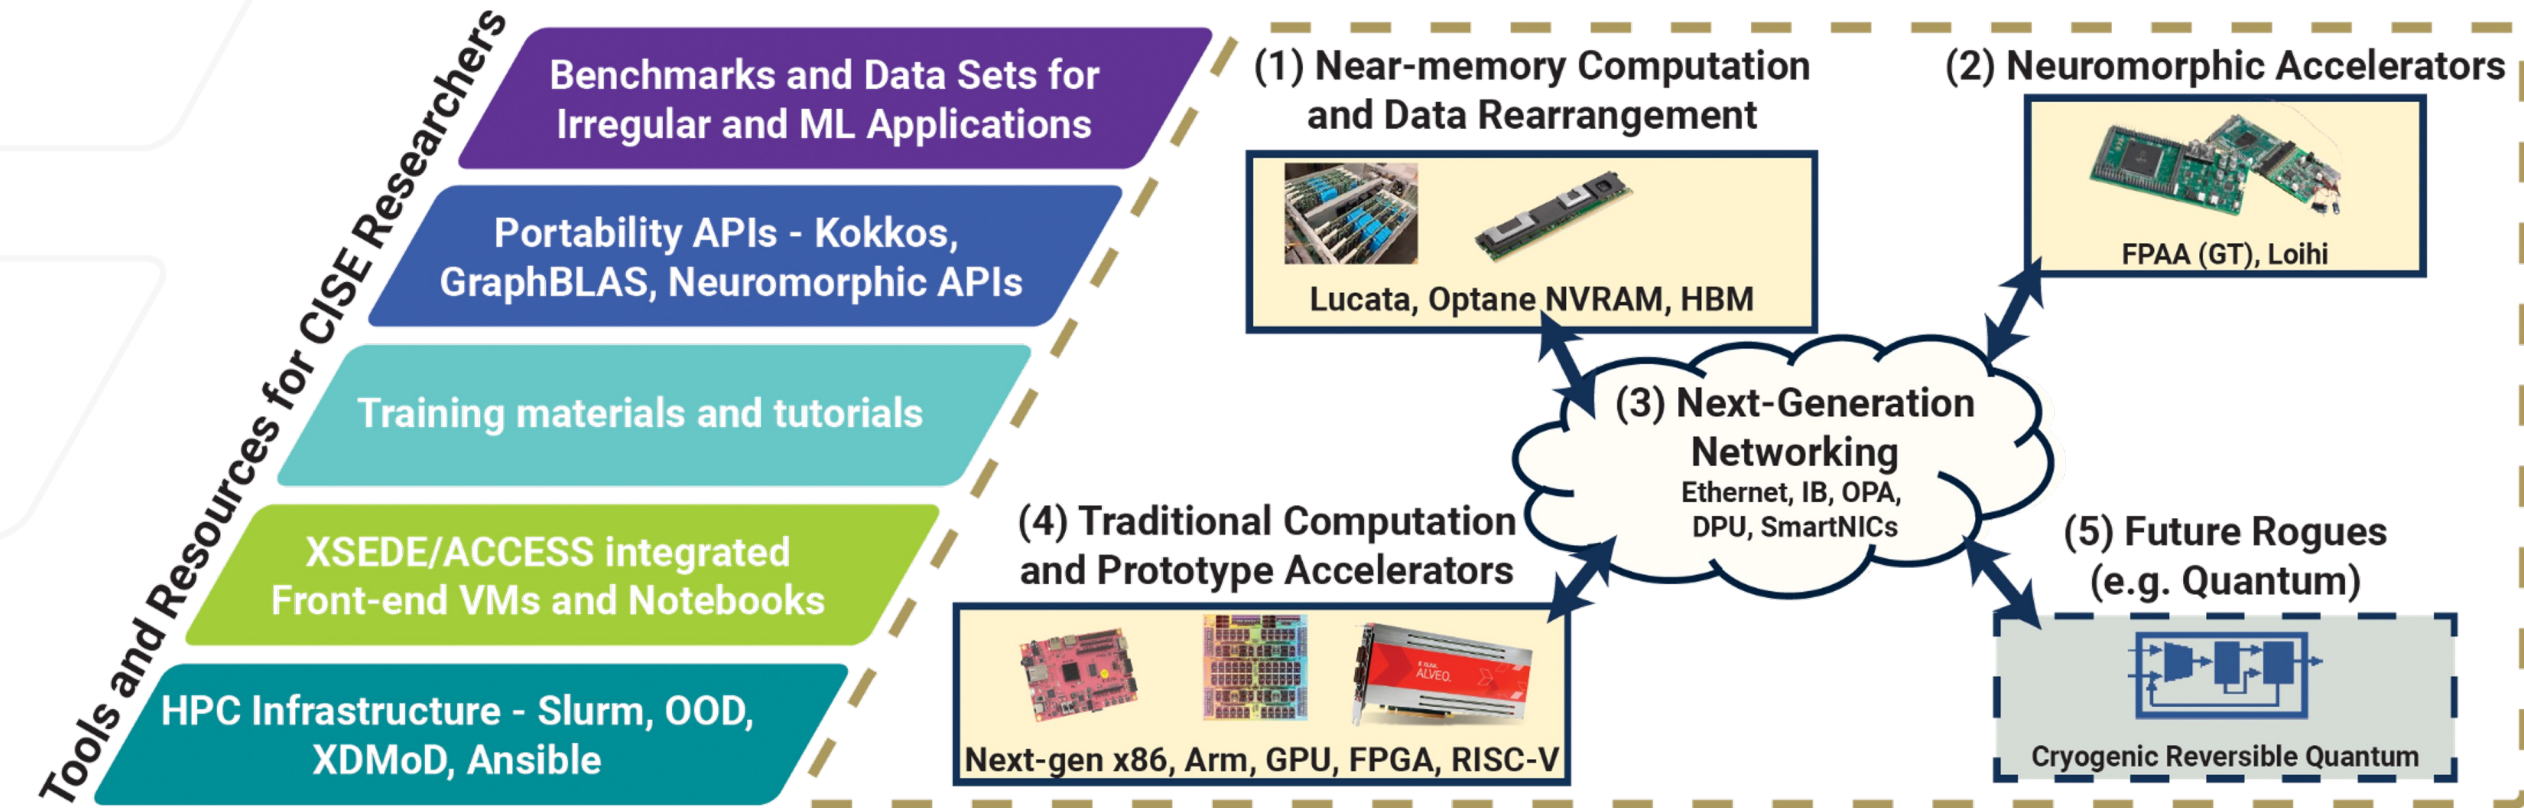
The Rogues Gallery an NSF funded post-Moore testbed for CISE researchers and the community.

CNS-2016701, \$1.3M over 3 years

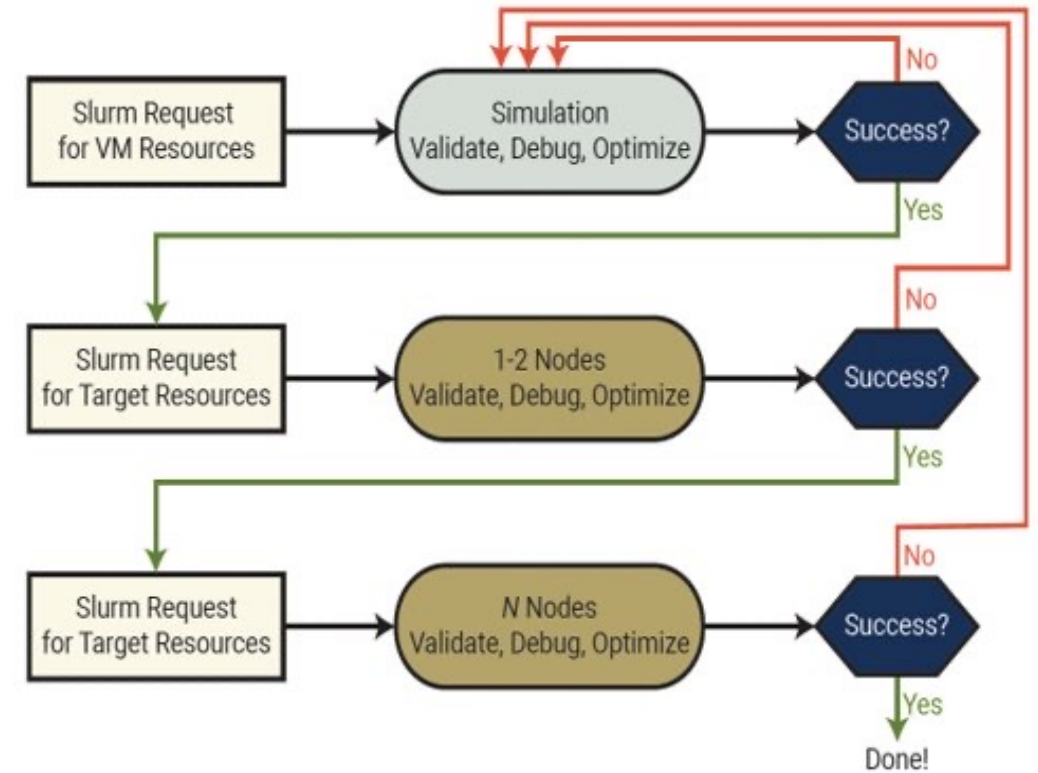
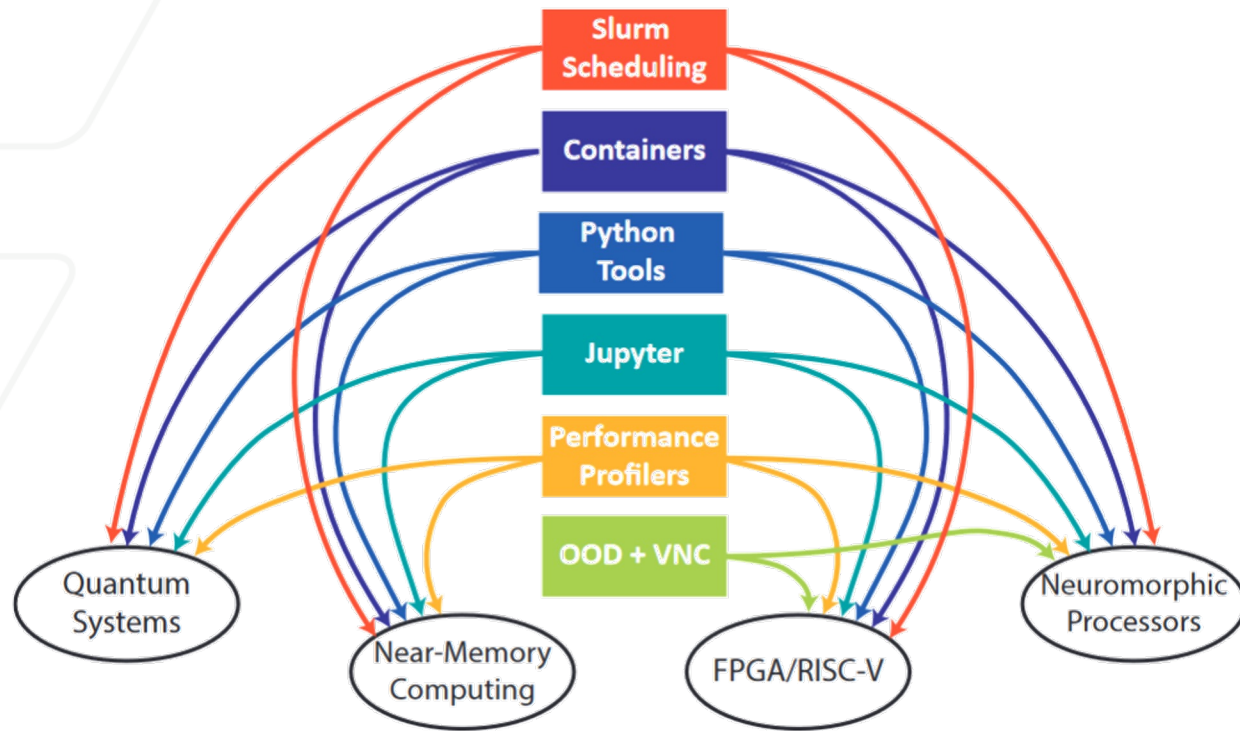
- Supports deploying:
 - Rack-scale Lucata Pathfinder 32 node system
 - Neuromorphic accelerators
 - Smart networking and 5G equipment
 - Backend infrastructure

This grant focuses on **community engagement and post-Moore training**

The Rogues Gallery



Managing Workflows for Novel Architectures



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

The VIP Program at Georgia Tech

- Vertically Integrated Projects (VIP) Program: late 1990s
 - Engage more students in research
- Teams are
 - Multidisciplinary (e.g. CS, ECE, physics)
 - Span student classifications (sophomore -> graduate)
 - Large-scale (>10 students)
- Projects are meant to be long-term
 - 3 semesters, 6 credits total (1, 2, 3 credits/semester)
 - Increasing leadership and technical ownership
- GT hosts more than 80 teams
 - 40 institutions world-wide form the VIP consortium



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

Future Computing with the Rogues Gallery VIP



- Started in Fall 2019
 - Drs. Jeffrey Young and Jason Riedy
- Vision focused on hardware on hand
 - Flagship “Emu Chick” for near-memory compute
 - Prototype FPGAs
 - Emerging HPC hardware (A64FX, Skylake, Epyc)
- Initial team comprised of 9 students
 - Big surprise was ask to explore quantum hw/sw
-> new subteam created accordingly!
- Today we have 35 students/semester + a waitlist



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

Today's Future Computing VIP

- Near-memory subteam
 - Focus on high-level frameworks (Cilk, GraphBLAS, etc.) and API calls
 - Micro benchmarks (pointer chase), macro benchmarks (BFS, hpcg)
- Quantum subteam
 - AI-driven circuit design
 - Train model based on Qiskit results -> simulate w/cuQuantum -> dispatch generated circuit
- Reconfigurable subteam
 - RISC-V architectures for specialized, efficient accelerators
 - Chisel/Chipyard + Vivado toolchain
- Neuromorphic subteam
 - Previously focused on SNNs for autonomous RC car navigating track
 - Extending to OpenAI Gym problems, targeting Lava + Loihi dispatch



Introduction



Project-Based
VIP



Extending to
Campus



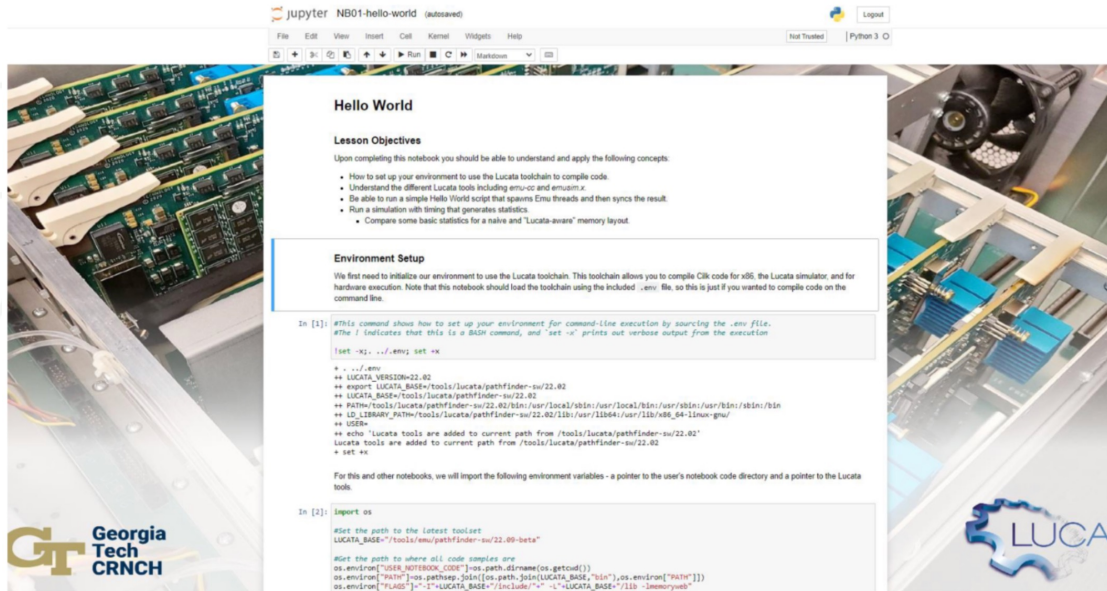
Challenges
Encountered



Summary

Critical Infrastructure for Success

- Presenting “funky” in a way that looks like “enterprise”
 - Slurm managed compute resources
 - Open OnDemand for interactive sessions
 - Lmod software stack
 - Apptainer containers for portability and simplicity
 - NFS mounts
- Abstracting complexity
 - 11 Slurm builds across 5 OSes!



Designing Curriculum to Foster Growth

- Student onboarding a big focus given vertical component of classes
 - Weekly how-to sessions focused on fundamentals, with increasing complexity
 - SSH -> git -> Slurm jobs -> Python environments -> Distributed workflows -> Containers
- Subteams develop documentation and examples to facilitate knowledge share
 - Hosted in central repo
- Weekly notebook entries to encapsulate work
 - Contributions, developments
 - Blockers
 - Links to documentation



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

Leveraging Knowledge to Provide PYNQ Capacity

- Curriculum change and supply change issues introduced new demand
 - CS3220: Processor Design required for all threads (100+ students/semester)
 - Additionally, ECE8893 FPGA Special Topics
 - Campus bookstore unable to stock sufficient FPGA development boards to meet need
- Used TechFee funds to procure 60 boards, set up Slurm-managed cluster
 - OOD VNC session on development VMs for Xilinx tools and overlay development, Jupyter for PYNQ



Introduction



Project-Based
VIP



Extending to
Campus

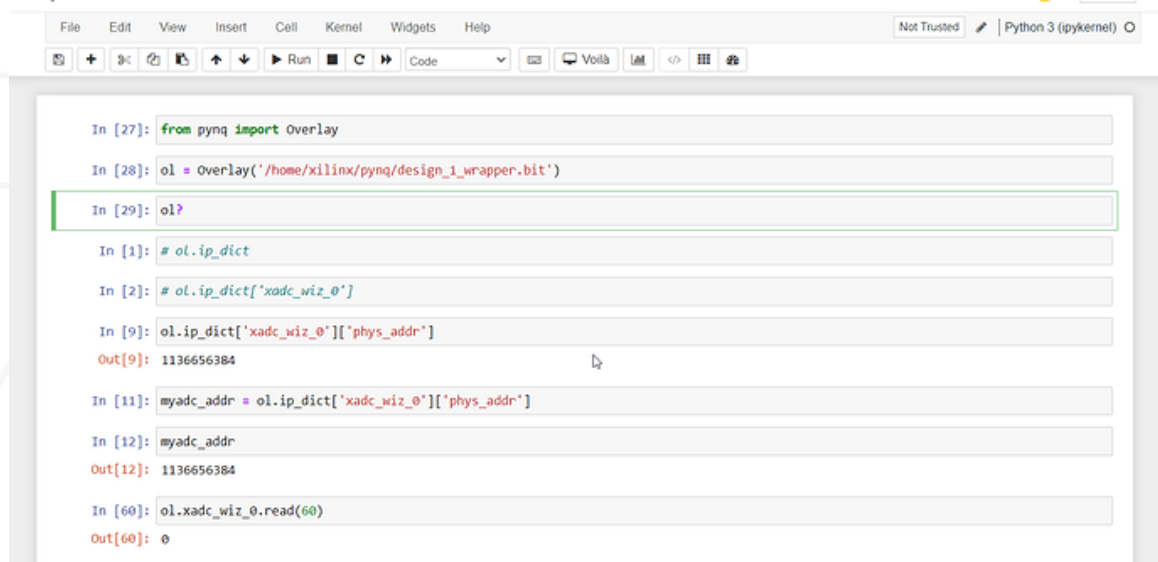


Challenges
Encountered



Summary

Hacking to a Functional Solution



```
In [27]: from pynq import Overlay
In [28]: ol = Overlay('/home/xilinx/pynq/design_1_wrapper.bit')
In [29]: ol?

In [1]: # ol.ip_dict
In [2]: # ol.ip_dict['xadc_wiz_0']
In [9]: ol.ip_dict['xadc_wiz_0']['phys_addr']
Out[9]: 1136656384

In [11]: myadc_addr = ol.ip_dict['xadc_wiz_0']['phys_addr']
In [12]: myadc_addr
Out[12]: 1136656384

In [60]: ol.xadc_wiz_0.read(60)
Out[60]: 0
```

- PYNQ boards are inherently designed for running python via Jupyter
 - Direct login as root
 - Notebooks, overlays expected to be in root-owned path
- Open OnDemand follows typical Jupyter job, with minor tweaks
 - /etc/sudoers includes a line allowing launch of jupyter lab with privilege
 - Prolog symlinks user's home directory into expected location

Scaling out to Central Infrastructure

- PACE resources (Phoenix and ICE) provide more capacity and capability, so as problems outgrow RG, extend workflows accordingly
 - Xilinx libraries on ICE for FPGA development
 - Chisel/Chipyard workflows on Phoenix
 - CuQuantum containers and support



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

Students Like Familiarity

- Like to work on laptops
 - Need encouragement to adopt cluster hardware
- Expectations like root access
 - Learning Python virtual environments, etc.
- Many of same challenges in regular clusters



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

Emerging Companies bring Growing Pains

- Emu chick -> Lucata Pathfinder transition
 - New architecture, scale, APIs
- Constantly evolving libraries
 - Often, semester to semester the standard would change and require time to rework existing code
- Sharing time with vendor engineers
 - Sharing prototype/development and “production” environments caused headaches
 - Lack of privilege for students meant pokes for support



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

...Including Disappearing

- Lucata recently announced that the company was being dissolved
 - No more engineering support for an *extremely* exotic system
- Taking the opportunity to encourage final publication of work/experience (PEARC?)
- Looking to the next project
 - A64FX, SPR-Max, GH100 all share near-memory concept
 - Can look to existing workflows on new architectures (with better support)



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary

In Summary

- For 5 years, we have supported Future Computing VIP
 - Seen significant growth
- We've extended efforts for other core curriculum
 - 120+ students/semester on PYNQ cluster - \$10k/semester in savings
- Continuing to add support for more exotic HW at GT!



Introduction



Project-Based
VIP



Extending to
Campus



Challenges
Encountered



Summary