

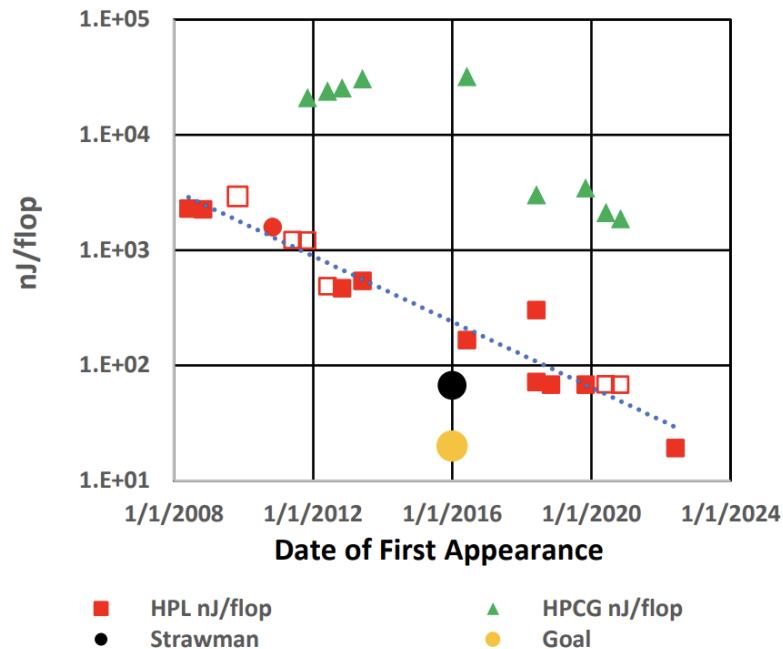
Novel Architecture Infrastructure for Addressing Future HPC Data Challenges

Jeffrey Young, PhD · Rogues Gallery Director, Principal Research Scientist · Partnership for Advanced Computing Environments

Durham University SciComp 2025 Seminar
February 20th, 2025

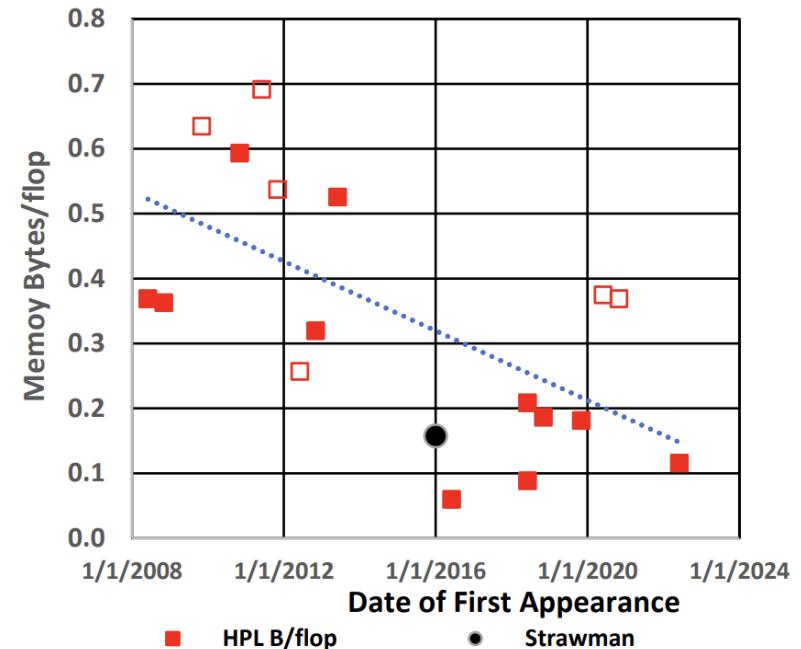
* Many thanks to Aaron Jezghani, SRS, PACE and Sara Karamati and Rich Vuduc for slide content

Exascale Data Movement Trends



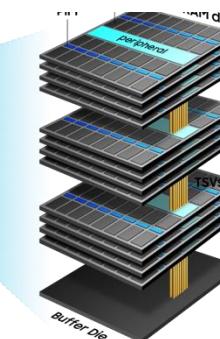
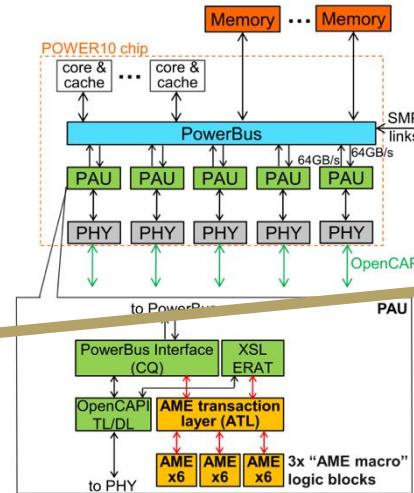
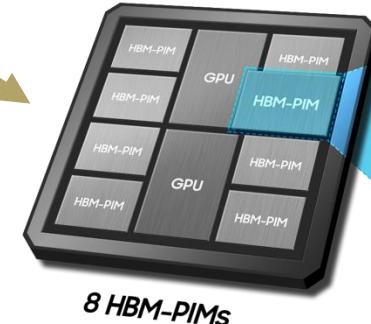
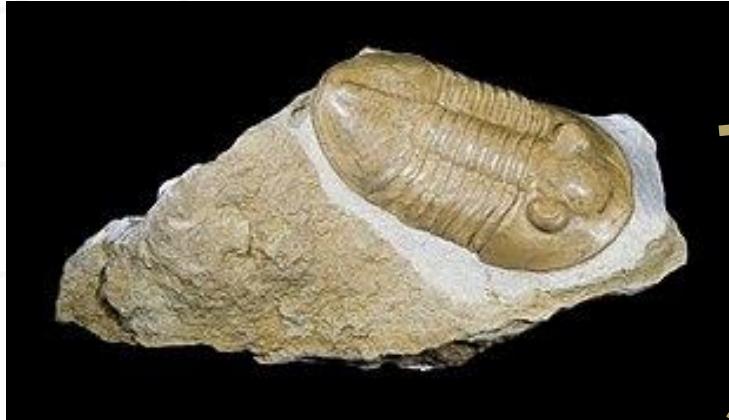
Energy/FLOP decreasing roughly matches 2008 projections

Decreasing Byte/FLOP ratio decrease indicates our current systems are trending in the wrong direction!



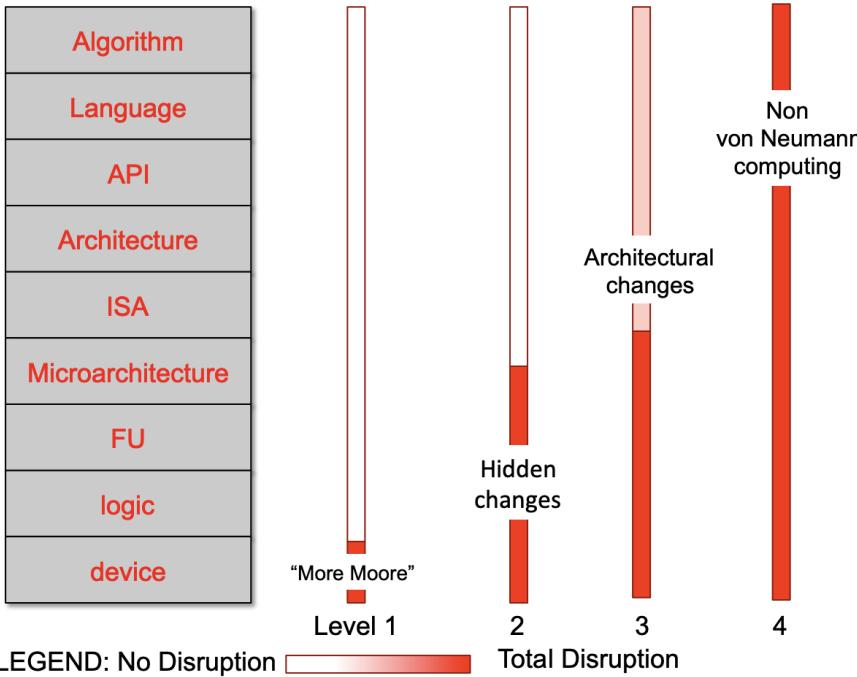
Clearly, we need to think differently about how we do data movement in the post-Moore and post-Exascale era!

Further Motivation



The “Cambrian Explosion” of new accelerators has led to many domain specific accelerators – not just in AI but also for data movement!

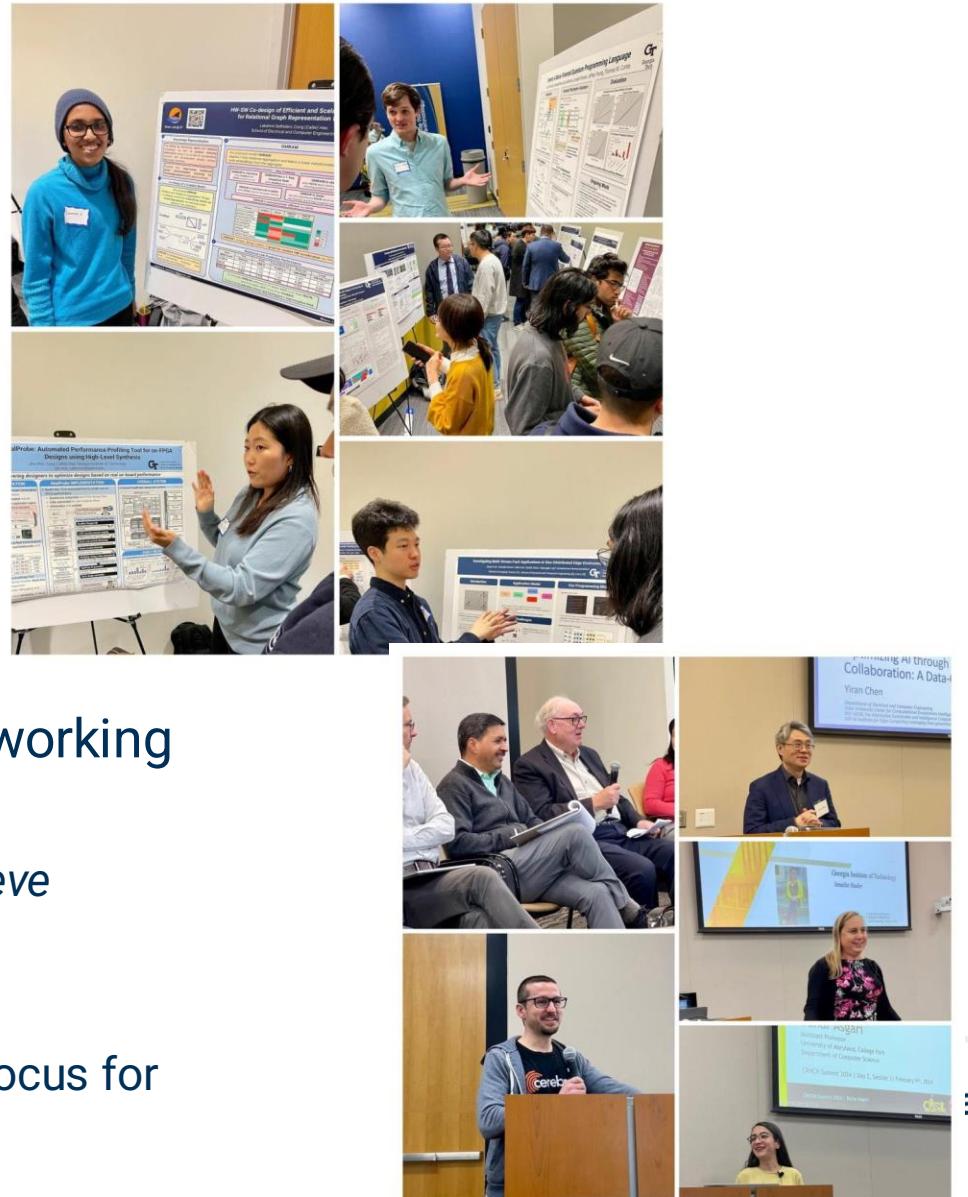
Center for Research into Novel Computing Hierarchies (CRNCH)



CRNCH was founded in 2017 to bring together researchers working all across the technology stack.

- *Both evolutionary and revolutionary improvements are needed to achieve performance in the post-Moore era*

The CRNCH Rogues Gallery testbed was created as a central point of focus for novel prototypes



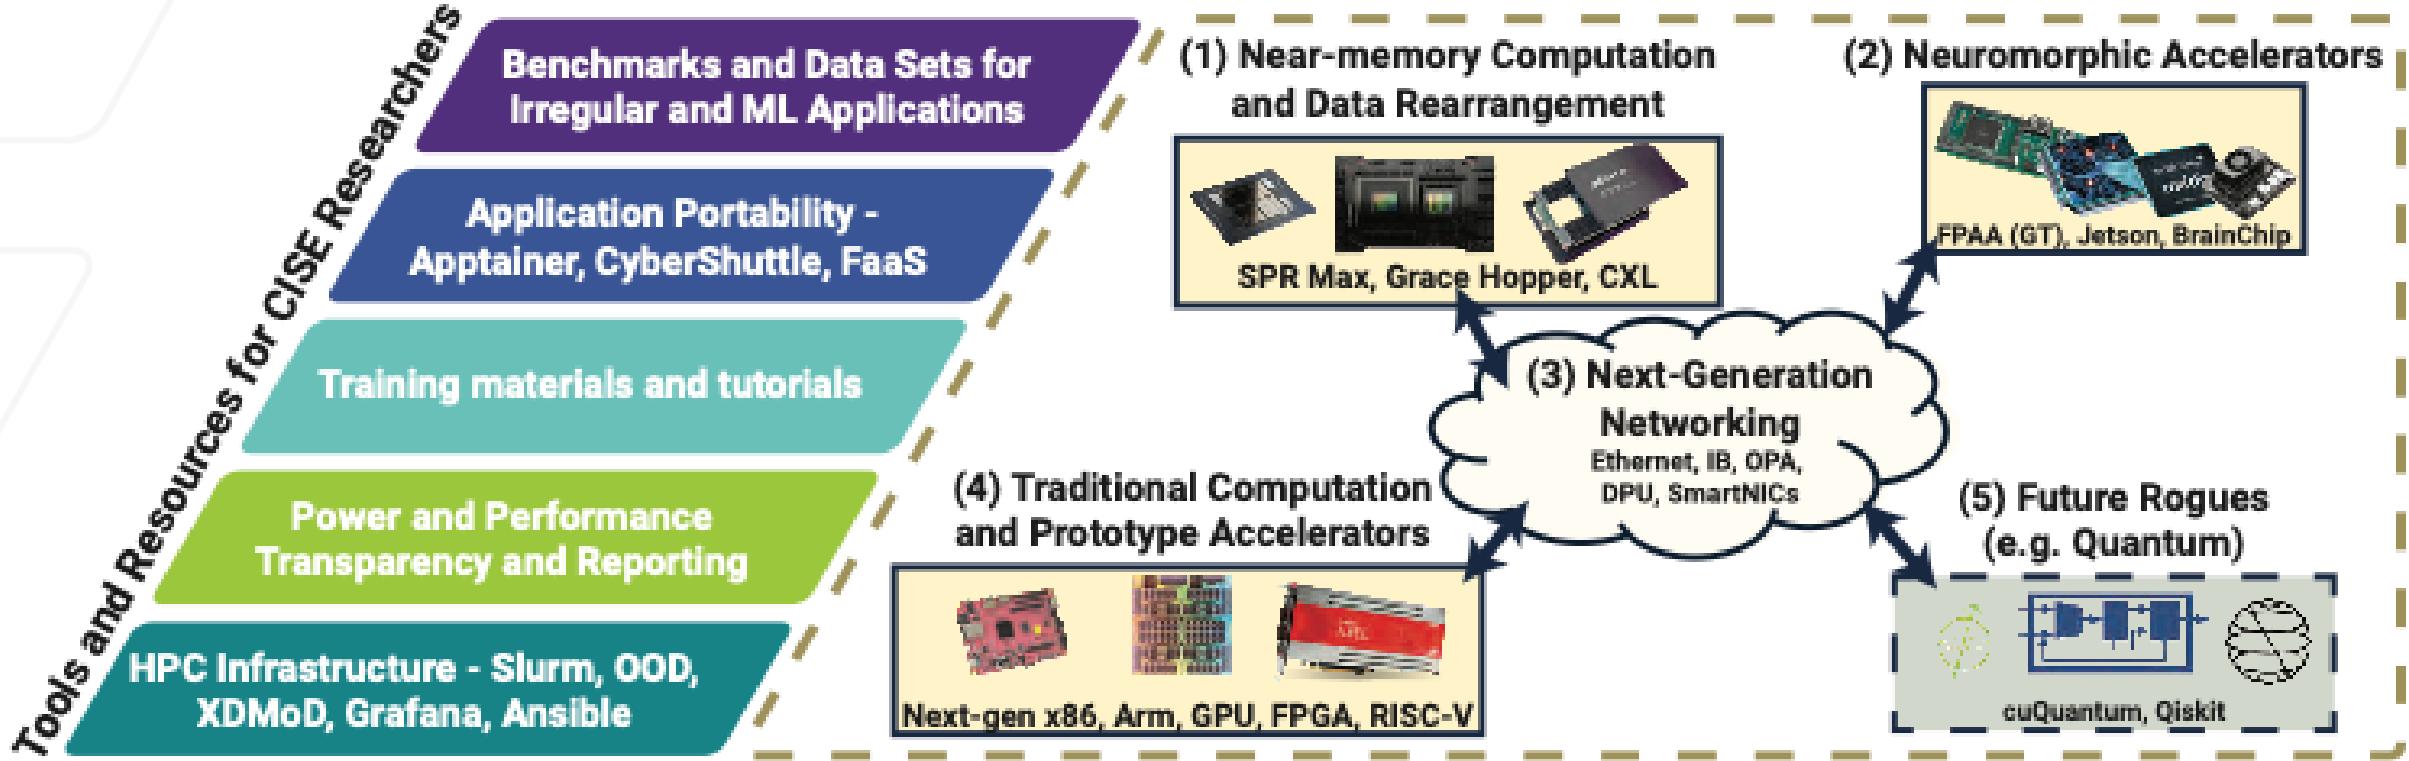
The Rogues Gallery



Rogues Gallery is an NSF funded post-Moore testbed for CISE researchers and the community.

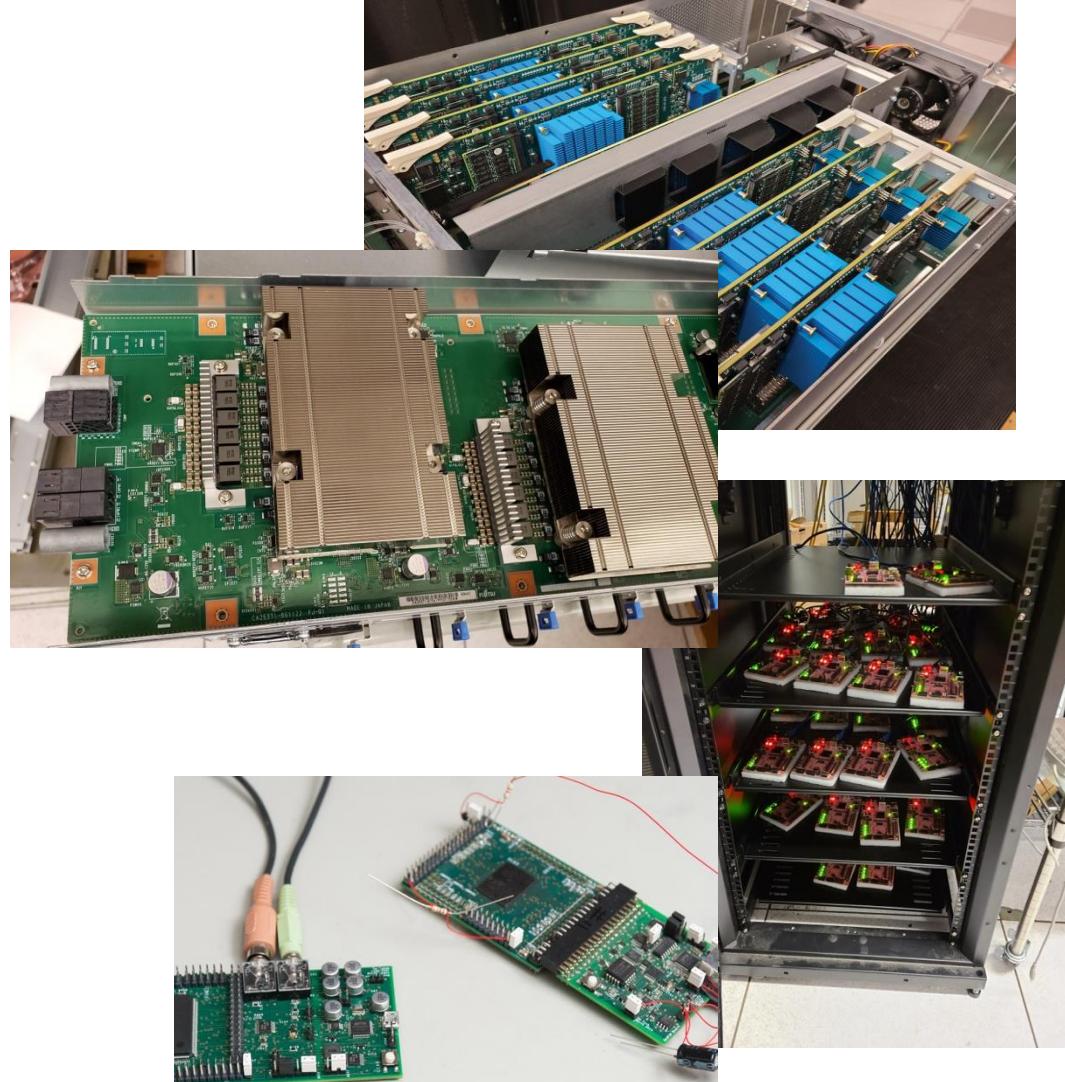
- The Gallery contains 40 servers and 40+ development boards – Intel CLX, SKL, ICX; AMD/NVIDIA GPUs; Arm; RISC-V; Xilinx
- ***Extreme heterogeneity*** with GPU, FPGAs, FPAAs, Optane Memory, InfiniBand, OmniPath, and Ethernet networking
- This grant focuses on ***community engagement and post-Moore training***

The Rogues Gallery



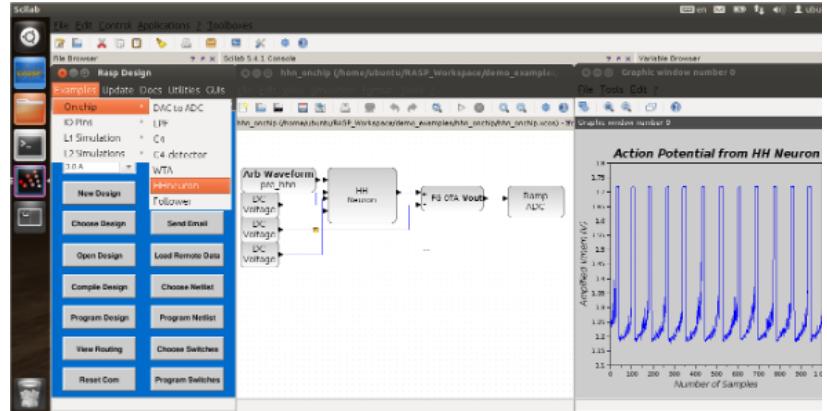
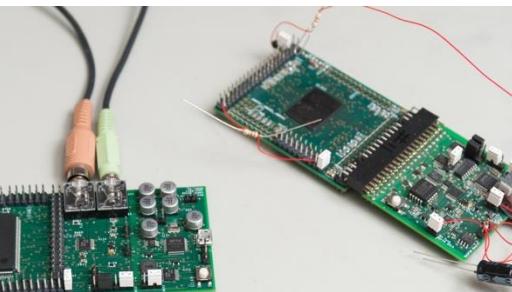
Rogues Gallery Highlights

- First deployment of A64FX with Open OnDemand, later adopted by RIKEN for Fugaku
- Largest public instance of Lucata Pathfinder - #211 on 2021 Graph500 and #46 on GreenGraph500 rankings
- Support for over 180 researchers and 60-70 external users across multiple areas
- Support for 80-130 students each year with Pynq boards and similar infrastructure

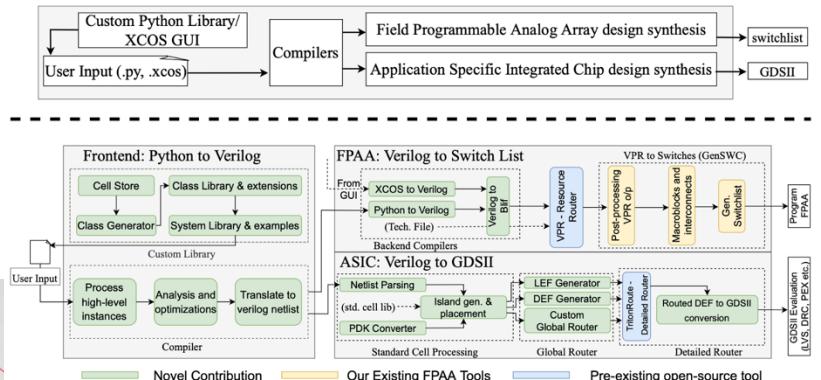


Unique Rogues Gallery Capabilities

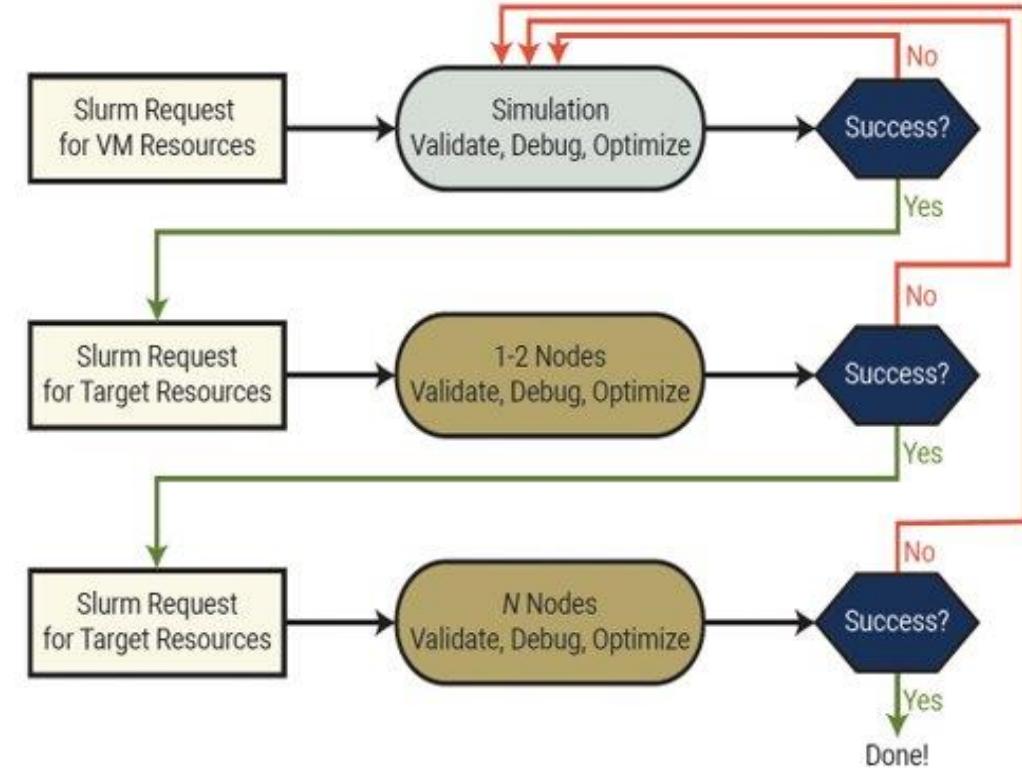
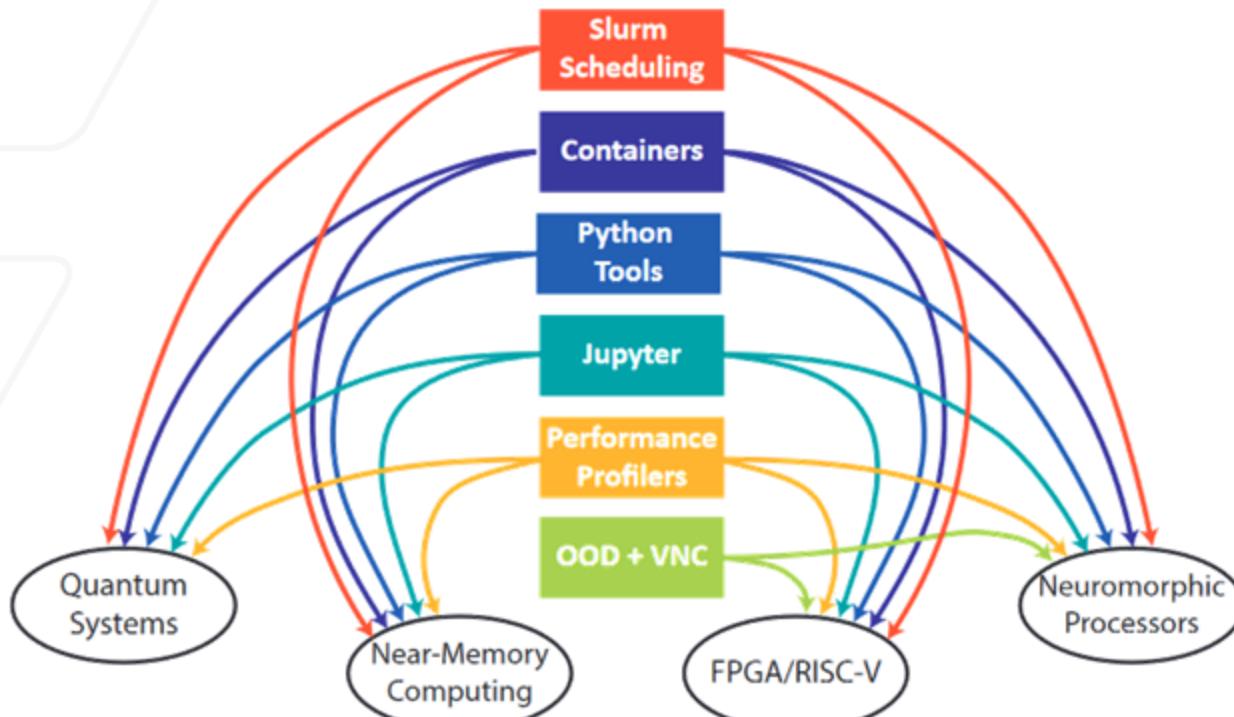
- Support for investigating graph analytics accelerators, neuromorphic hardware, and smart networking
- Integration with local educational mission
 - Provides valuable testing of novel architecture
 - Ex: FPAA Python workflow, PYNQ cluster
- Growing interactions with software development best practices
 - CI/CD, continuous benchmarking



Next generation tools for FPAA development:
graphical with Xcos interface (above), full open-source analog FPAA workflow (below)

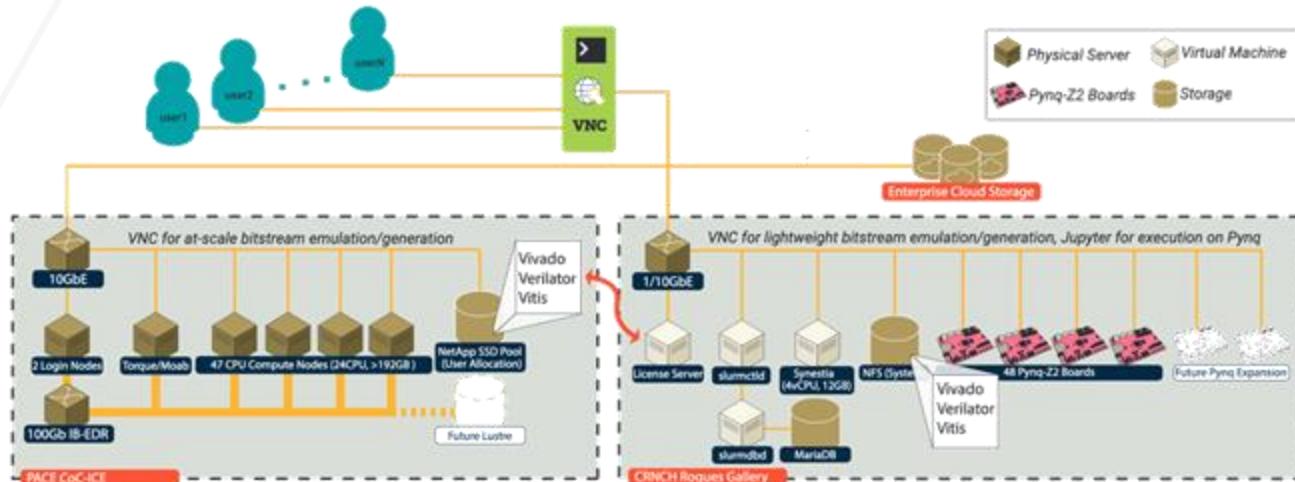


Managing Workflows for Novel Architectures



Slurm and Open OnDemand

- Slurm scheduling across most resources
 - Interesting challenges with “edge device” integration
 - Running 24.11 with 16 different builds (x86, Arm, NVIDIA, RISC-V, etc.)
- Open OnDemand support to ease onboarding



The screenshot shows the Georgia Tech OnDemand interface. The top navigation bar includes links for Georgia Tech, Files, Jobs, Clusters, Interactive Apps, My Interactive Sessions, and Dev. The main content area displays the CRNCH (Center for Research into Novel Computing Hierarchies) logo and the "ROGUES GALLERY". Below the logo, it says: "OnDemand provides an integrated, single access point for all of your HPC resources." A "Message of the Day" section follows, and a list of available job submission options is provided. At the bottom, it says "powered by OPEN OnDemand".

CI/CD support on the Rogues Gallery

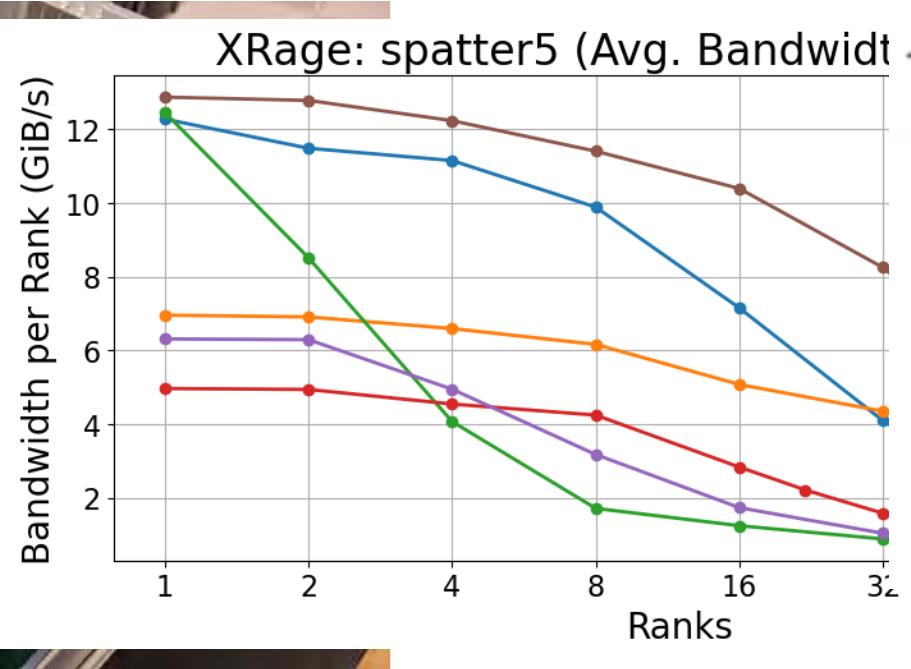
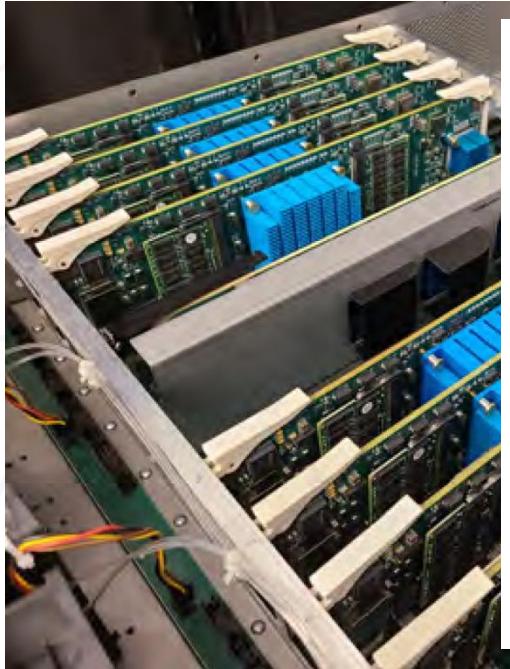
- Part of supporting novel architectures includes supporting *effective software development*
 - CRNCH RG provides Slurm-oriented infrastructure to schedule runners for novel architectures (RISC-V, Arm, Intel GPU, FPGA)
 - Dedicated VM and secure setup process allows for persistent runners



254 workflow runs			
		Event ▾	Status ▾
		Branch ▾	Actor ▾
✓ Build	Build #215: Scheduled	main	11 hours ago 7m 47s
✓ Build	Build #214: Scheduled	main	yesterday 7m 38s

Learn More: <https://gt-crnch-rg.readthedocs.io/en/main/general/ci-runners.html>

Case Studies of How We Use the RG Testbed



Lucata Pathfinder – Move the Compute to the Data

Spatter - Simulating Next-Generation Memory Systems

SmartNICs – What are they Good For?

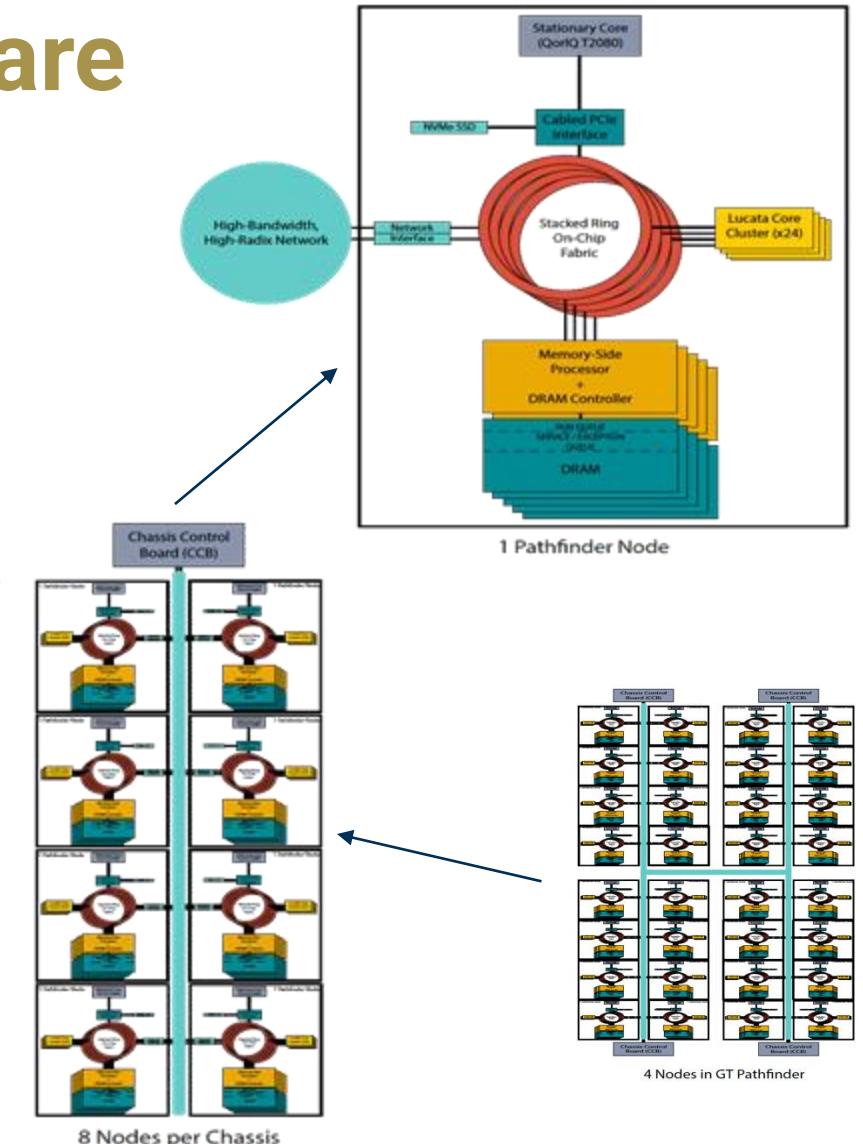
Case Study 1 – Move the Compute to the Data



- The Lucata Pathfinder is the next iteration of Emu Chick prototype
- Design motivated by scaling issues for problems hindered by data access patterns
 - E.g. graphs, sparse matrix algebra
- Rather than migrating large swaths of data across network, hardware threads can move to the memory
 - Demonstrated scaling efficiency with benchmarks such as BFS, hpcg vs. x86

Pathfinder - Reconfigurable Hardware

- Each chassis contains:
 - 1 Chassis Control Board (PowerPC E6500) for node/network management
 - 8 compute nodes, which include:
 - 1 stationary compute core (PowerPC E6500)
 - 24 Lucata compute cores (Proprietary IP)
 - Stacked ring-on-chip fabric to interface to network/DRAM/compute
 - High-bandwidth, high-radix network for inter-node/chassis communication
- Systems can be configured as single or multi-chassis
 - Defines network and available Lucata compute cores for running applications
 - Scale hardware according to the problem size
- GT Pathfinder consists of 4 chassis



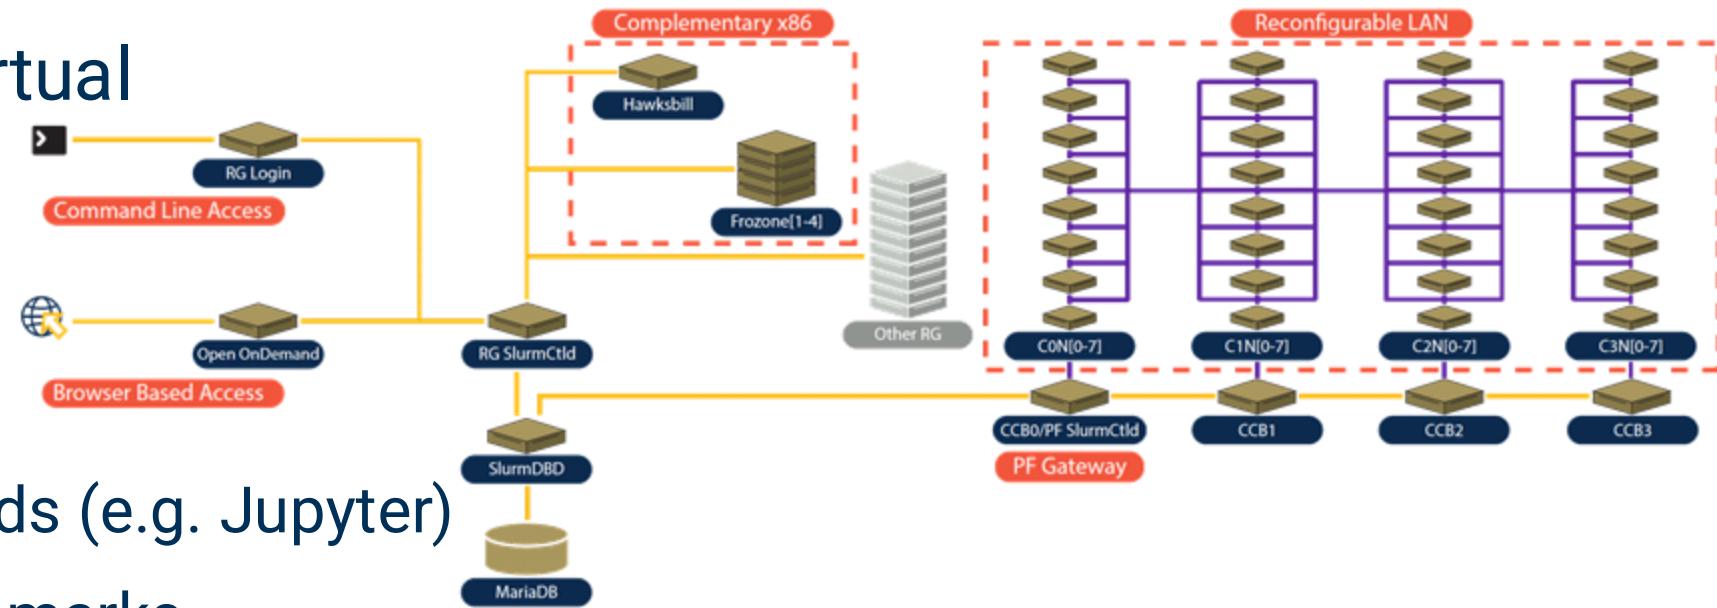
High-Level Libraries/APIs for Code

- Lucata-provided software environment
 - Custom adaptation of Yocto Linux kernel and base utilities
 - Configuration tools to dynamically modify chassis, network
 - Performance monitoring libraries to check efficiency, energy
 - Custom API for thread management and task execution
 - Robust simulation capabilities to validate before deployment
- Open-Cilk for parallel programming
 - Open-source platform for task-parallel code development
 - Leverages lightweight backend ABIs for target dispatch
 - Minimal effort to port code from x86 to Pathfinder

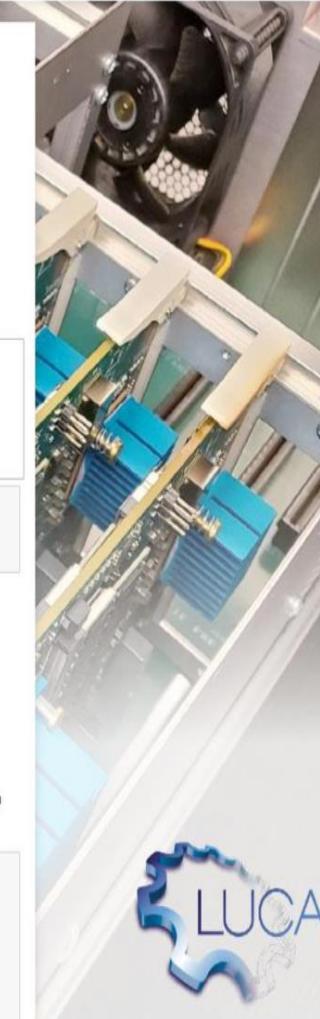
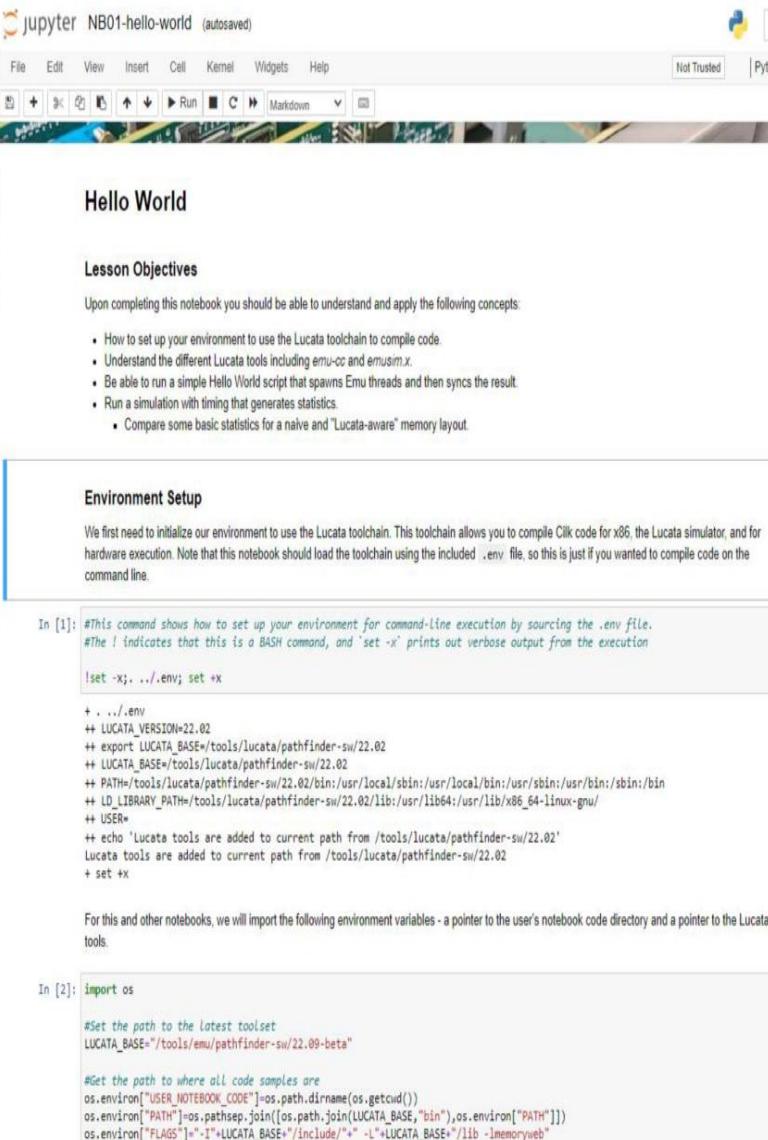


Accessing the Pathfinder and Related Systems

- Pathfinder access managed via Slurm (and calendar schedule)
 - Federated cluster alongside other Rogues and infrastructure nodes
 - Single entry point, single job database
- x86 physical and virtual machines for development and simulation
 - Host visual frontends (e.g. Jupyter)
 - Comparative benchmarks



Community Tutorials to Engage Researchers



Hello World

Lesson Objectives

Upon completing this notebook you should be able to understand and apply the following concepts:

- How to set up your environment to use the Lucata toolchain to compile code.
- Understand the different Lucata tools including `emu-cc` and `emusim`.
- Be able to run a simple Hello World script that spawns Emu threads and then syncs the result.
- Run a simulation with timing that generates statistics.
 - Compare some basic statistics for a naïve and "Lucata-aware" memory layout.

Environment Setup

We first need to initialize our environment to use the Lucata toolchain. This toolchain allows you to compile Cilk code for x86, the Lucata simulator, and for hardware execution. Note that this notebook should load the toolchain using the included `.env` file, so this is just if you wanted to compile code on the command line.

```
In [1]: #This command shows how to set up your environment for command-line execution by sourcing the .env file.
#The ! indicates that this is a BASH command, and 'set -x' prints out verbose output from the execution
!set -x; ./../.env; set +x

+ ./../.env
++ LUCATA_VERSION=22.02
++ export LUCATA_BASE=/tools/lucata/pathfinder-sw/22.02
++ LUCATA_BASE=/tools/lucata/pathfinder-sw/22.02
++ PATH=/tools/lucata/pathfinder-sw/22.02/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
++ LD_LIBRARY_PATH=/tools/lucata/pathfinder-sw/22.02/lib:/usr/lib64:/usr/lib/x86_64-linux-gnu/
++ USER=
++ echo 'Lucata tools are added to current path from /tools/lucata/pathfinder-sw/22.02'
Lucata tools are added to current path from /tools/lucata/pathfinder-sw/22.02
+ set +x

For this and other notebooks, we will import the following environment variables - a pointer to the user's notebook code directory and a pointer to the Lucata tools.

In [2]: import os

#Set the path to the latest toolset
LUCATA_BASE="/tools/emu/pathfinder-sw/22.09-beta"

#Get the path to where all code samples are
os.environ["USER_NOTEBOOK_CODE"] = os.path.dirname(os.getcwd())
os.environ["PATH"] = os.pathsep.join([os.path.join(LUCATA_BASE, "bin"), os.environ["PATH"]])
os.environ["FLAGS"] = "-I" + LUCATA_BASE + "/include/" + "-L" + LUCATA_BASE + "/lib -lmemoryweb"
```

Georgia Tech CRNCH

Tutorials covering Emu Chick/Lucata Pathfinder presented at numerous conferences

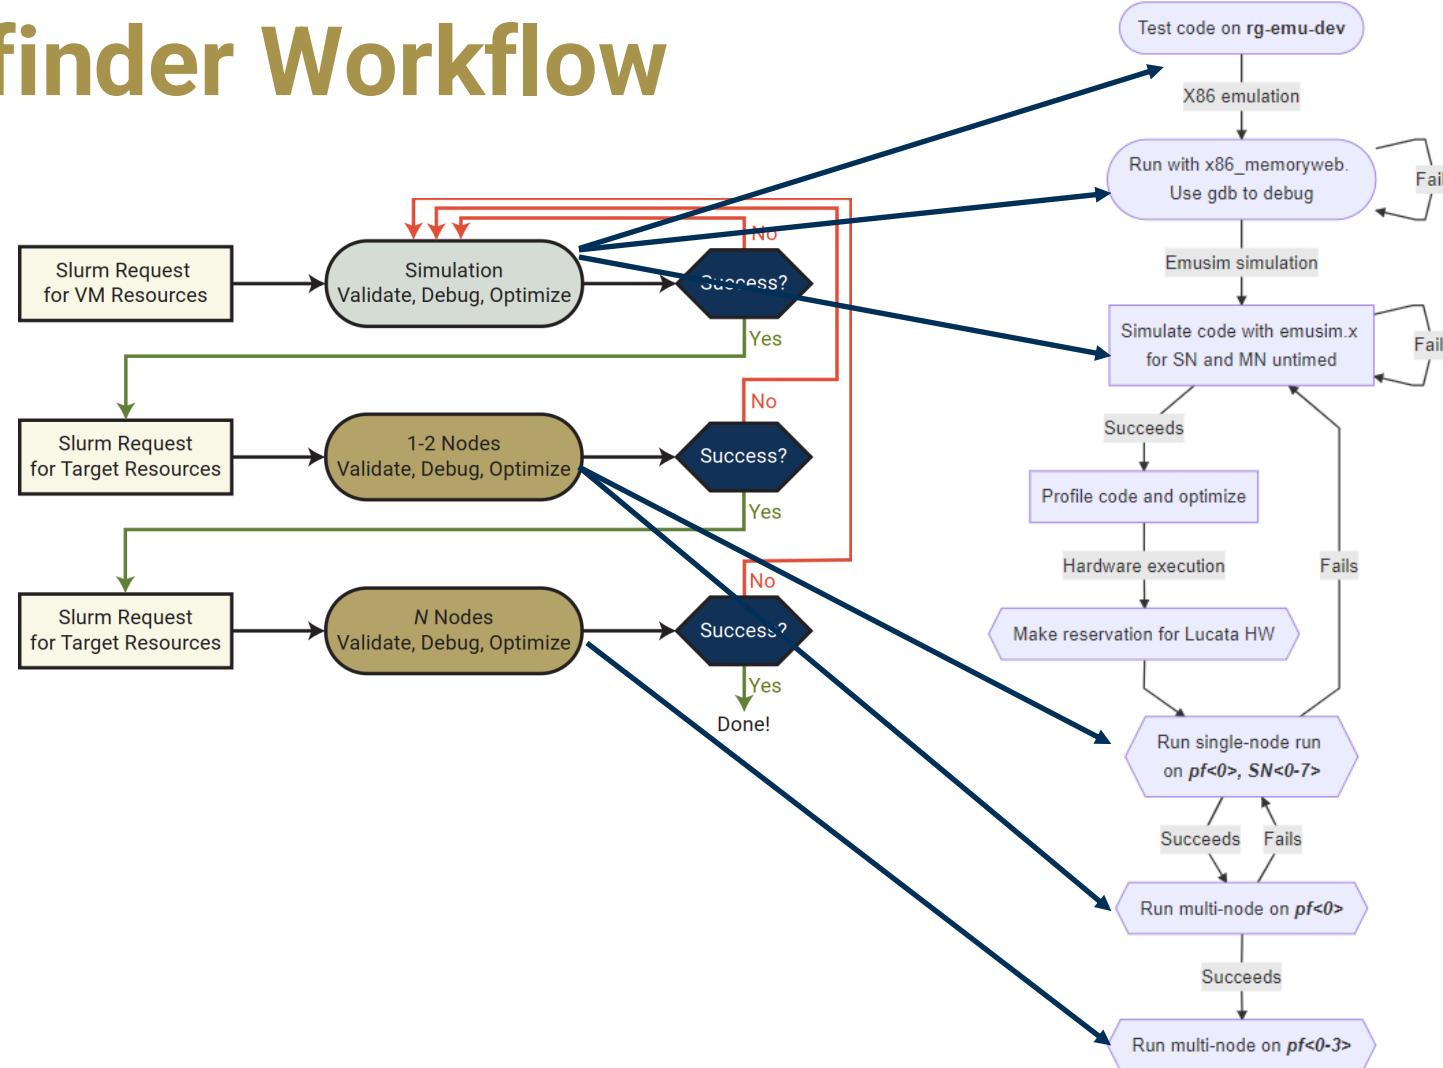
- General RG overview at ASPLOS19 and PEARC19
- Dedicated tutorial at PEARC21 and HPEC22
- Iteratively developed content and workflows, such as Jupyter-based approach for HPEC22

Lessons learned for the platform and tutorials in general:

- User engagement depends on user interface and prerequisite knowledge
- Conferences with fewer cross-listed and hybrid options increased attendance and participation
- Attendees may have HPC experience, but still need more focused examples for new architectures

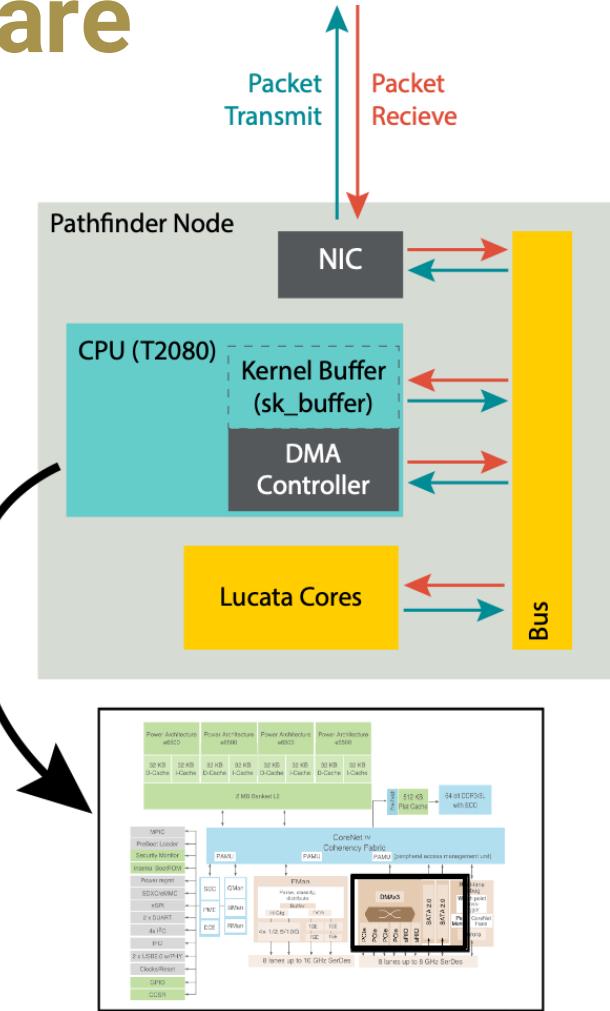
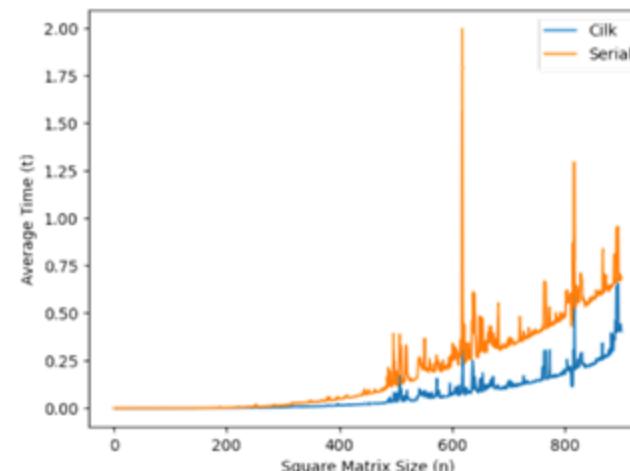
Full Lucata Pathfinder Workflow

Working with the Lucata Pathfinder still requires a complex workflow, but we can break it down into small, concrete steps for users

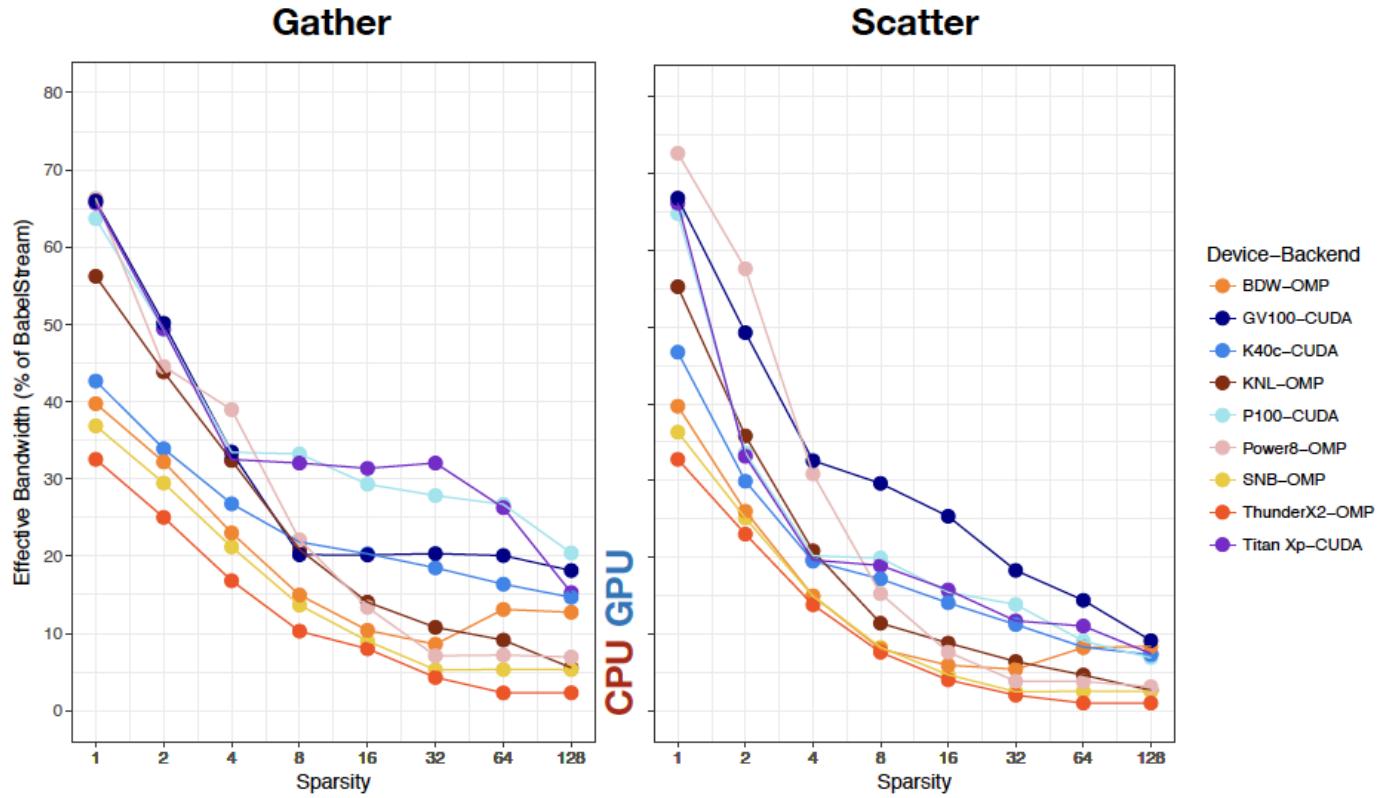


Extending CS Curriculum to Novel Hardware

- Future Computing with the Rogues Gallery VIP
 - Near-memory subteam focuses on the Pathfinder architecture for parallel and distributed computing
- Explore benchmarks on x86 and Pathfinder systems
 - hpcg to assess scaling performance for SpMV
 - BFS to demonstrate graph capabilities
 - Pointer Chase (pChase) to compare memory latency
- DMA driver prototype for Lucata compute cores
 - Project option for CS3210: Design of Operating Systems



Case Study 2 – Evaluating Next-generation Architecture Designs for Memory Systems



See Kevin Sheridan, et. al, "A Workflow for the Synthesis of Irregular Memory Access Microbenchmarks." In *Proceedings of the International Symposium on Memory Systems*, pp. 219-234. 2024.

Motivation

	Memory subsystem						Floating Point		Non-FP
	L1	L2	L3	DRAM	DRAM BW	Mem Latency	DP FLOPs	Vectorization	
Flag 3D Ale	Green	Yellow	Green	Green	Green	Green	2.50%	7.10%	97.50%
PartiSN 42 groups	Red	Green	Yellow	Red	Red	Red	26.20%	90.40%	73.80%
Jayenne DDMC Hohlraum	Yellow	Green	Green	Red	Green	Yellow	14.30%	0.20%	85.70%
xRAGE Shaped Charge	Red	Green	Yellow	Yellow	Red	Yellow	6.50%	14.00%	93.50%
Application 1	Yellow	Green	Green	Yellow	Red	Yellow	7.80%	19.20%	92.20%
Application 2	Yellow	Green	Green	Yellow	Red	Yellow	8.10%	17.60%	91.90%

Hardware Bottlenecks for LANL HPC codes [1]

However, we still have little understanding of how future memory accelerators might affect codes of interest because:

- Finding appropriate regions of interest (ROI) for the memory system is challenging to infer even with tools like SimPoints and LoopPoint
- Some applications that we might want to benchmark can't be fully shared to extract meaningful traces or ROIs

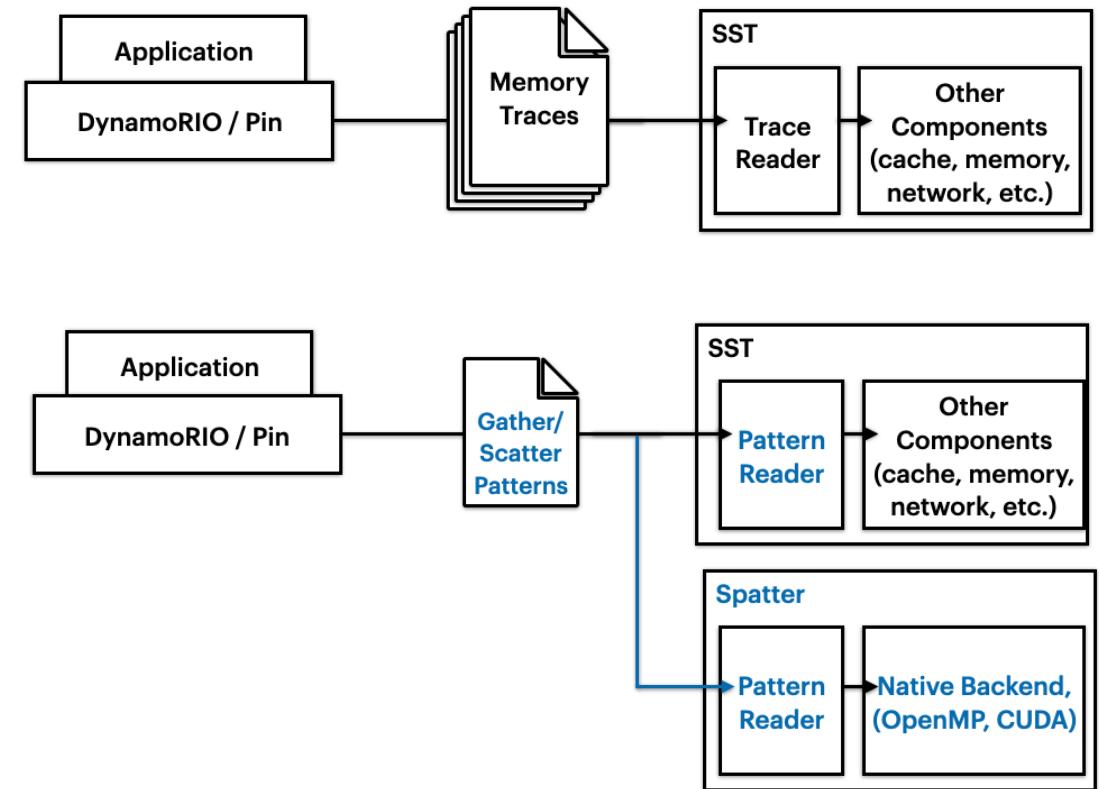
Our ideal workflow would allow us to 1) capture relevant memory accesses from real-world applications, 2) benchmark them on real systems, and 3) simulate new hardware models

[1] G. Shipman, et al., Assessing the Memory Wall in Complex Codes, MCHPC 2022 doi: 10.1109/MCHPC56545.2022.00009

Using Codesign to Design Better Memory Systems

Our ideal tool workflow would allow us to:

- Capture relevant memory accesses from real-world applications
- Run these patterns on bleeding edge systems
- Use the ***same patterns*** to simulate the performance of future near-memory accelerators



Spatter - Version 1.0

The basis of Spatter is two kernels; one for gather and another for scatter.

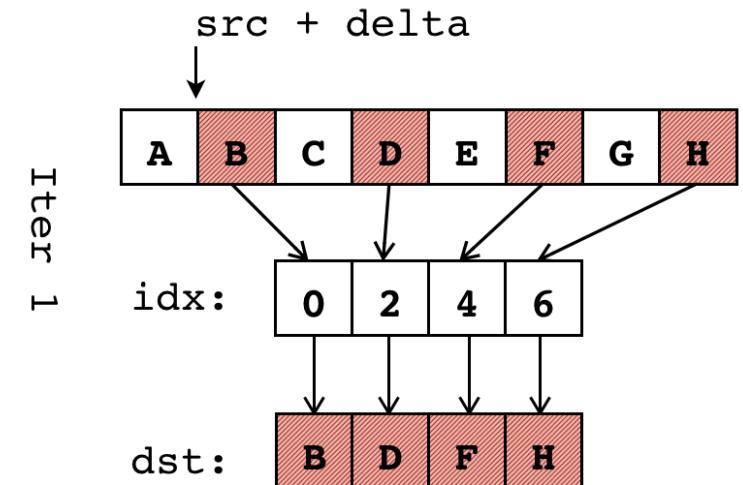
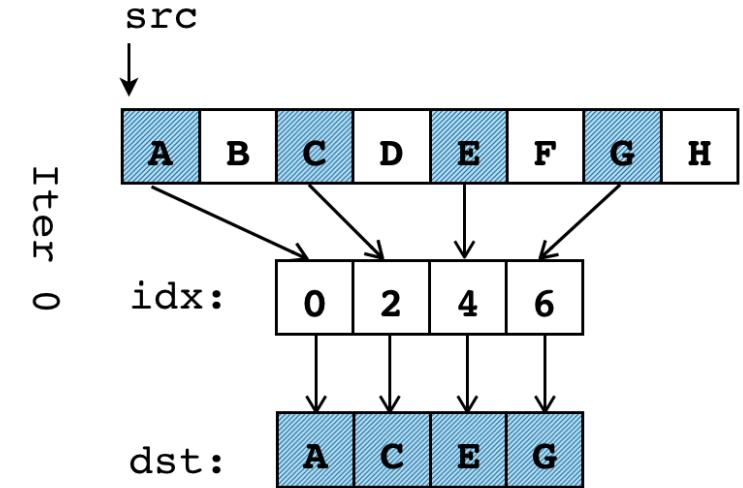
Gather kernel:

```
for i in 0..N:  
    reg = gather(src + delta*i, idx)
```

Scatter kernel:

```
for i in 0..N:  
    scatter(dst + delta*i, idx, reg)
```

The delta and the pattern in idx specify the *memory access pattern*.



GS Patterns

Sliding window approach is used to track non-trivial memory accesses within an application trace or region of interest (ROI)

Multiple filters are used to keep the most relevant access patterns for the final pattern output

GS Patterns codebase at
https://github.com/lanl/gs_patterns

	maddr1	maddr2	...	maddrN	iaddr
ScatterWindow =	0x0480	0x0488	...	0x1488]	0x0001
	0x1780	0x1788	...	0x0783	0x0003

	0x9580	0x5588	...	0x3581	0x1019

Filter #	Description
1	Index distances are only -1, 0, and/or 1
2	No symbol
3	Not in top 10 window appearance counts
4	Less than 1024 instances
5	Less than 6 unique index distances and less than 50% out of bounds distances*

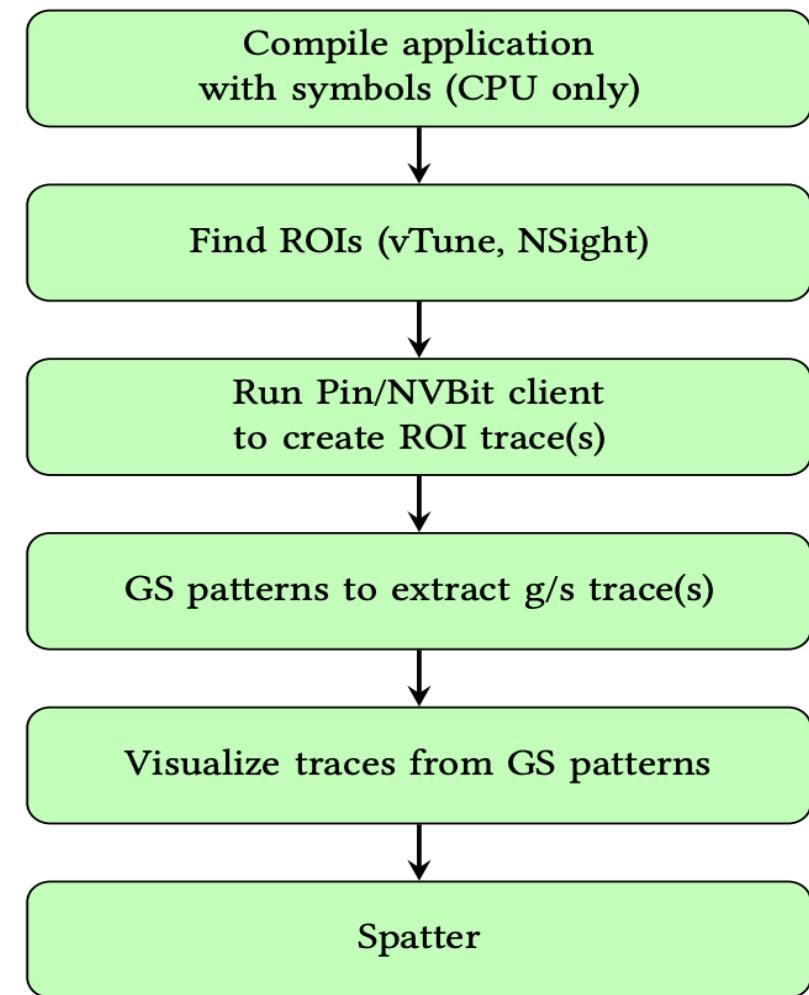
*Filter rules apply to each full memory access sequence. Sub-sequences are not removed. *Default out of bounds distances in $(-\infty, -513]$ or $[513, \infty)$.*

GS Patterns Workflow

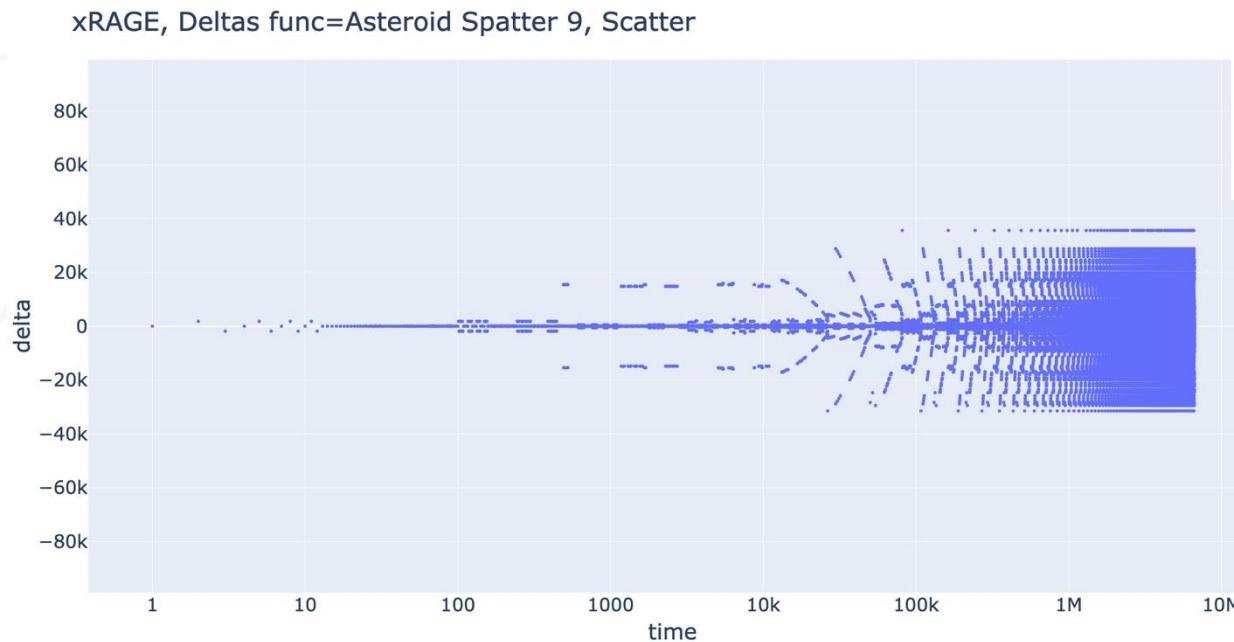
Further extensions to GS Patterns refactored code to use C++ and a plugin infrastructure for PinTool, NVBit

Currently ROI analysis speeds up the overall GS Patterns workflow but is a somewhat manual process to run and annotate codes

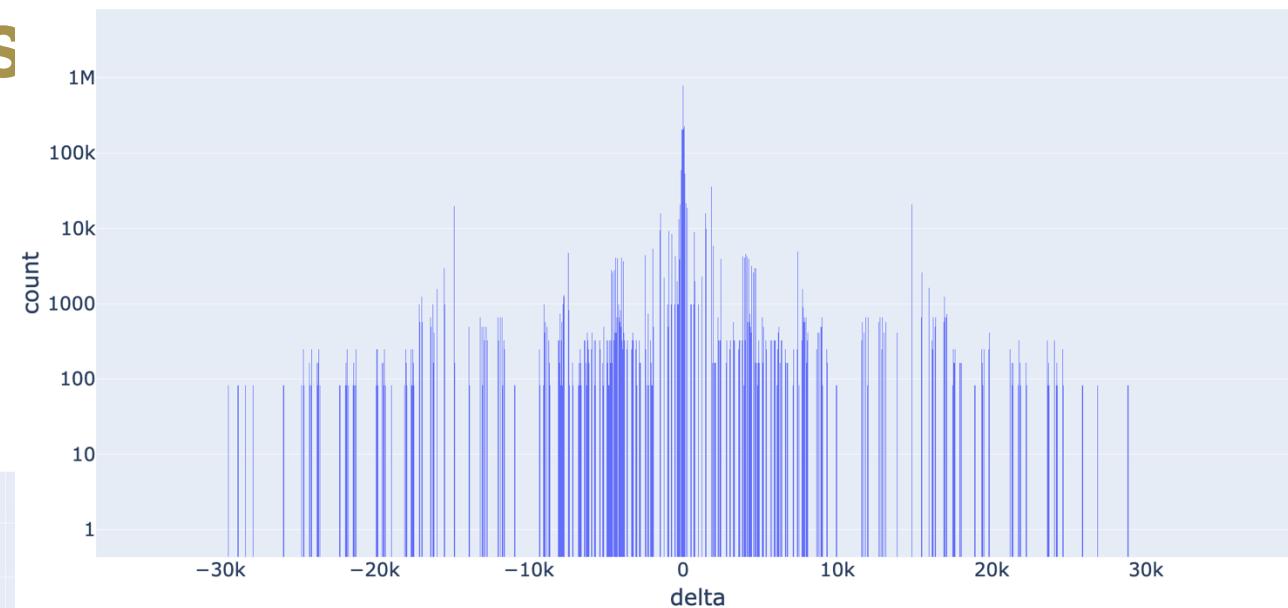
*Check out collected public patterns at
<https://github.com/hpcgarage/spatter-patterns>*



Pattern Visualization Tools



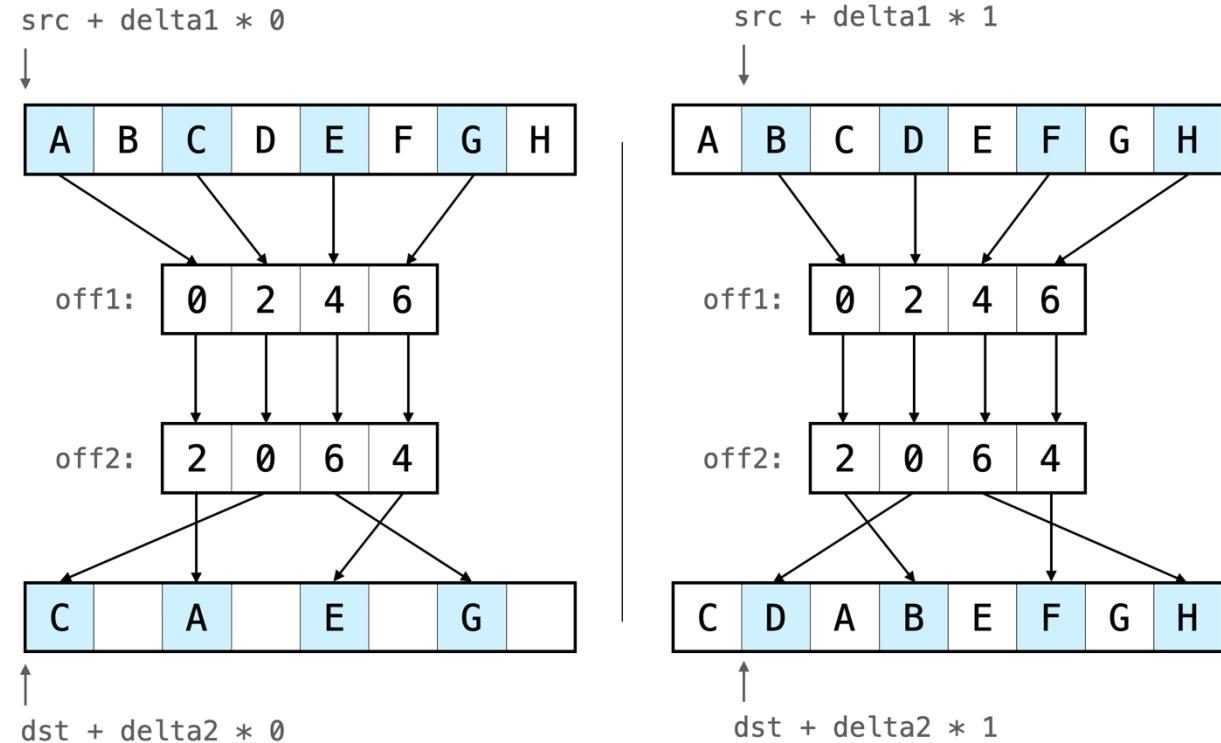
xRAGE, Deltas Histogram func=Asteroid Spatter 9, Scatter



The GS Patterns JSON output can be easily visualized in a variety of different ways...

Spatter 2.0

- Complete refactor of argument parsing, build system, and movement towards C++ design
- Addition of new kernels to better represent multiple levels of indirection - GatherScatter, MultiGather, MultiScatter
- Support for longer offset buffer lengths for improved application pattern representation
- MPI support for weak/strong scaling
- Support for atomics with scatter operations



GatherScatter Kernel Representation

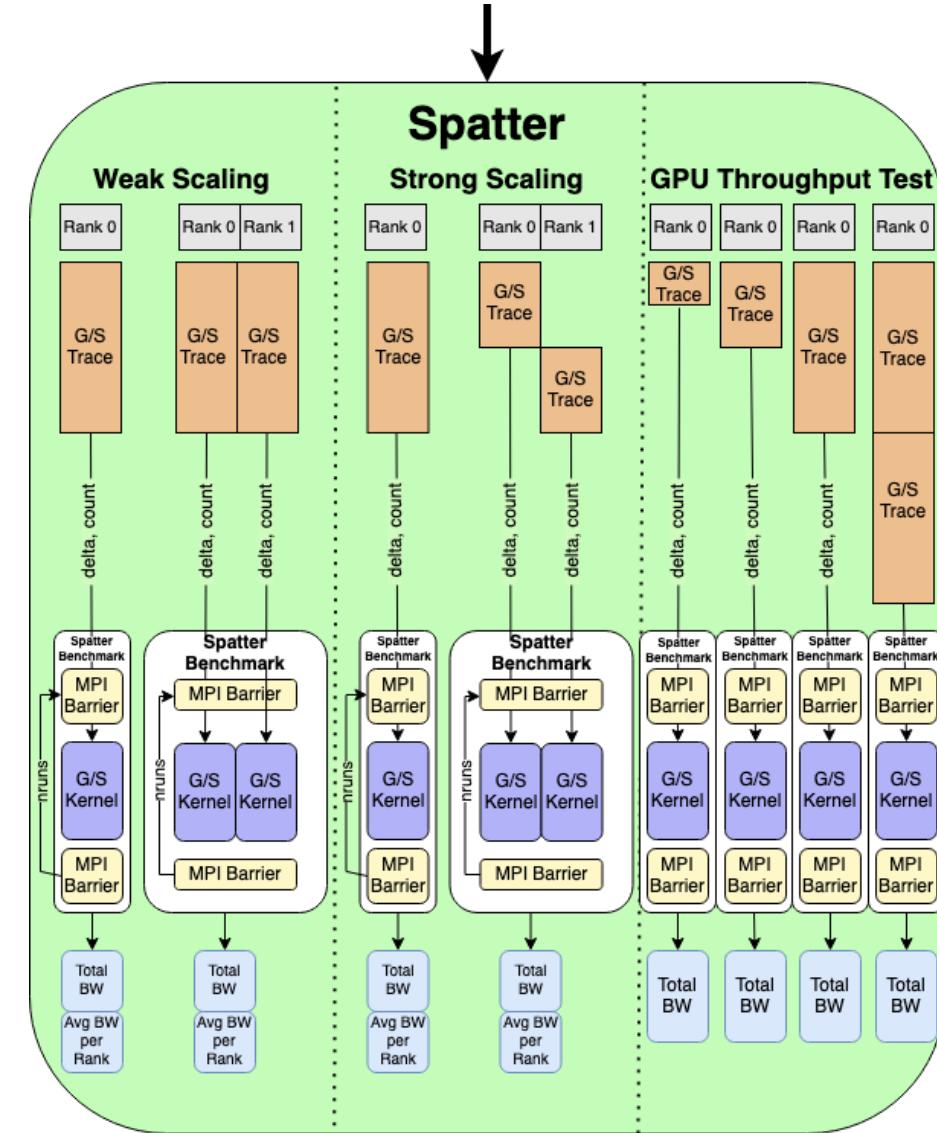
<https://github.com/hpcgarage/spatter>

Spatter 2.0 MPI Workflow

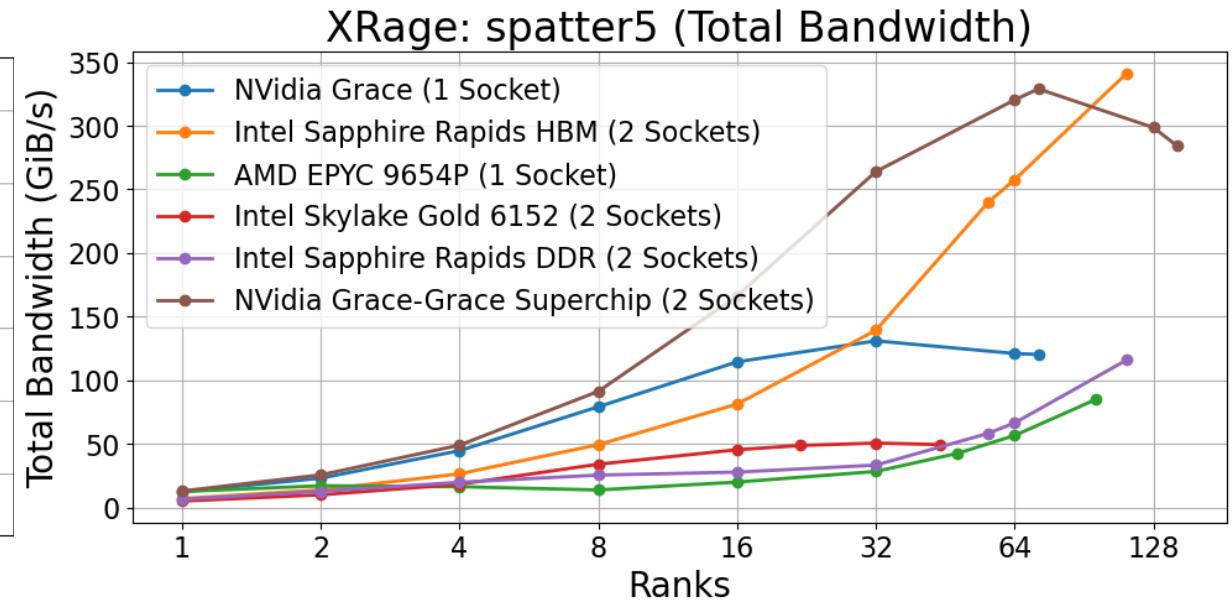
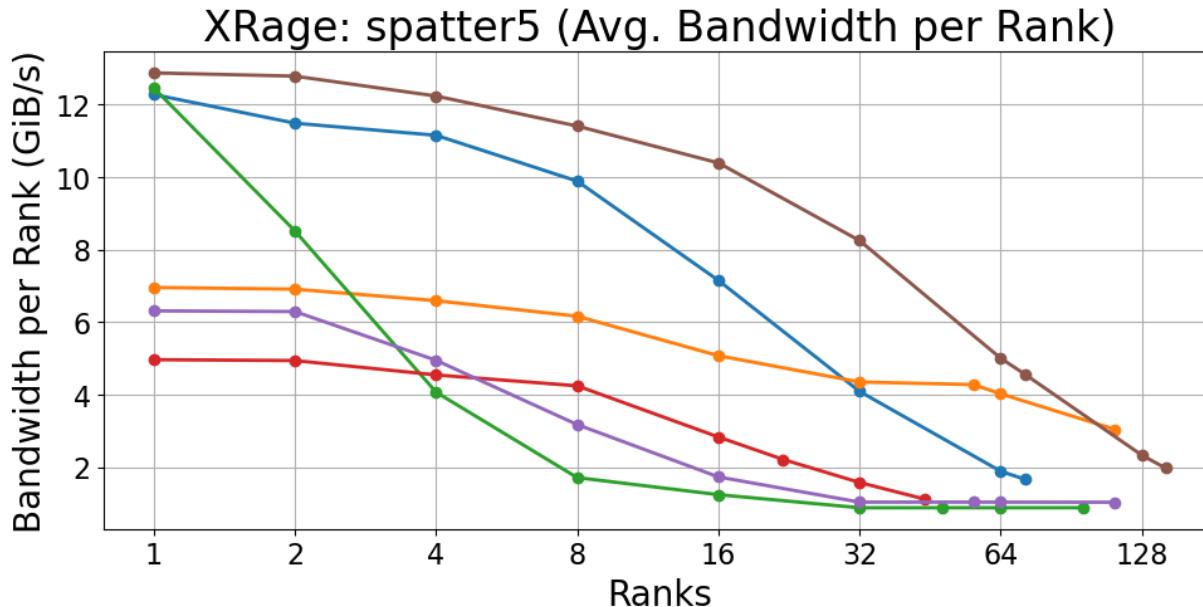
Weak Scaling – each MPI rank gets the same pattern and scaling scripts are used to sweep across N ranks

Strong Scaling – access patterns are partitioned across MPI ranks

GPU Throughput – patterns are truncated or expanded to vary amount of memory accesses and saturate GPU memory subsystem



Spatter 2.0 Results



Weak Scaling MPI Tests for Xrage Gather pattern

For more details, please see: Kevin Sheridan, et. al, "A Workflow for the Synthesis of Irregular Memory Access Microbenchmarks." In Proceedings of the International Symposium on Memory Systems, pp. 219-234. 2024.

Spatter Integration with SST



```
Running Spatter version 1.1
Compiler: GNU ver. 11.4.0
Backend: Serial
Aggregate Results? NO
```

```
Run Configurations
[ {'id': 0, 'kernel': 'scatter', 'pattern': [0, 24, 48, 72, 96, 120, 144
    'pattern-gather': [], 'pattern-scatter': [], 'delta': 0, 'delta-gath
    ...
{'id': 11, 'kernel': 'gather', 'pattern': [0, 24, 48, 72, 96, 120, 144
    'pattern-gather': [], 'pattern-scatter': [], 'delta': 1, 'delta-gath
```

config	bytes	time(s)	bw(MB/s)
0	73959168	0.0025168	29386.1
1	29593344	0.00101654	29111.8
2	21479040	0.000900899	23841.8
3	16384256	0.00103612	15813.1
4	12334080	0.000430775	28632.3
5	12334080	0.000427073	28880.5
6	12311808	0.000426331	28878.5
7	11265408	0.000443739	25387.5
8	9829632	0.000339094	28987.9
9	9829632	0.000783912	12539.2
10	9829632	0.000334772	29362.2
11	9250560	0.000322094	28720.1

Run Configurations

```
[ {'id': 0, 'kernel': 'scatter', 'pattern': [0, 24, 48, 72, 96, 120, 144, 168, 192, 216, 240, 264, 288, 312, 336, 360],
    'pattern-gather': [], 'pattern-scatter': [], 'delta': 0, 'delta-gather': 8, 'delta-scatter': 8, 'count': 577806, 'wrap': 1, 'threads': 1},
    ...
{'id': 11, 'kernel': 'gather', 'pattern': [0, 24, 48, 72, 96, 120, 144, 168, 192, 216, 240, 264, 288, 312, 336, 360],
    'pattern-gather': [], 'pattern-scatter': [], 'delta': 1, 'delta-gather': 8, 'delta-scatter': 8, 'count': 72270, 'wrap': 1, 'threads': 1} ]
```

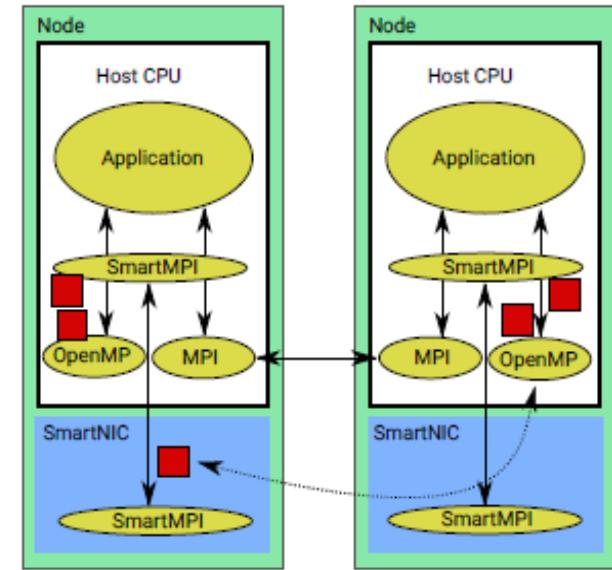
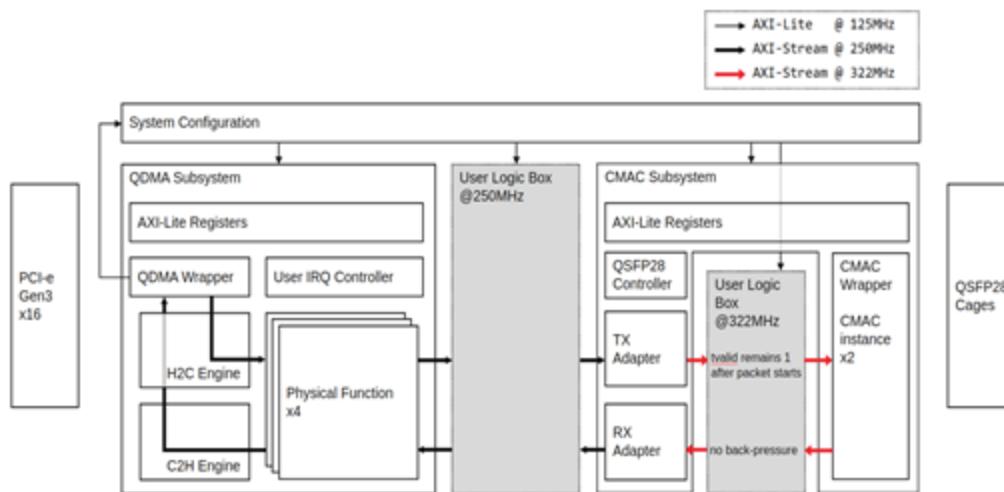
config	bytes	time(s)	bw(MB/s)	cycles	time(s)/cycles
0	73959168	0.0104005	7111.11	5489157	1.89474e-09
1	29593344	0.00410376	7211.27	2300269	1.78404e-09
2	21479040	0.00297854	7211.26	1669566	1.78402e-09
3	16384256	0.00230403	7111.12	3393131	6.79029e-10
4	12334080	0.00163812	7529.40	1259456	1.30066e-09
5	12334080	0.00154176	8000.01	2530289	6.09321e-10
6	12311808	0.00153897	8000.01	2525826	6.09296e-10
7	11265408	0.0015622	7211.25	826265	1.89067e-09
8	9829632	0.00136309	7211.30	2024452	6.73312e-10
9	9829632	0.00136309	7211.28	10223443	1.3333e-10
10	9829632	0.00136308	7211.32	1212079	1.12458e-09
11	9250560	0.00128279	7211.26	678285	1.89123e-09

```
Simulation is complete, simulated time: 44.5121 ms
```

Lulesh Output (i5-12400F Alder Lake vs. SST Miranda Spatterbench)

Larger application patterns may benefit from either parallel simulation sweeps or by using MPI-based tests

Case Study 3 – SmartNICs for HPC?



Example of task-based runtime work by Weinzierl, et al. “*Intelligent algorithms on intelligent networks—experiences and challenges using NVIDIA’s BlueField technology*”. Slides at <https://dpu.ornl.gov/>

What's a SmartNIC?

The basic building block of a distributed-memory cluster or supercomputer is a node.

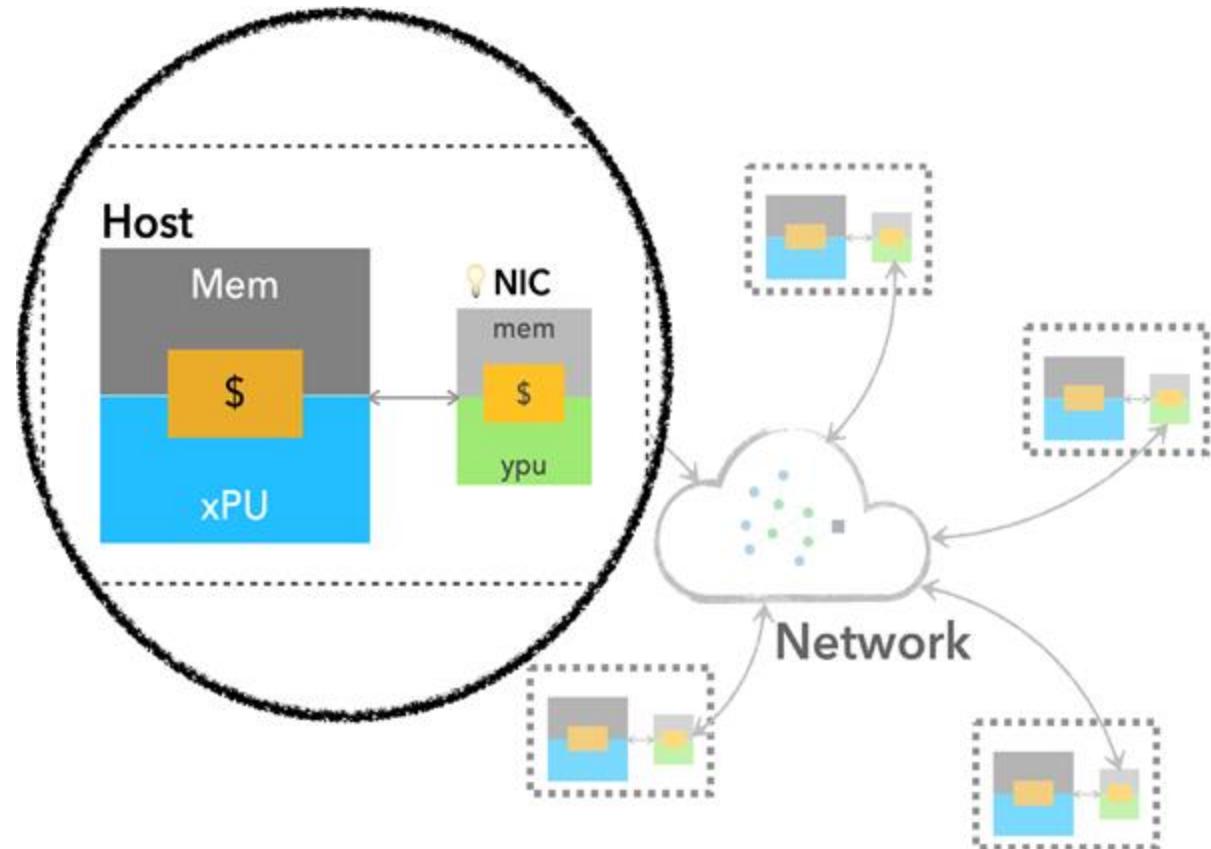
Each node includes a host, which is a processor (xPU) + memory hierarchy.

The host can communicate with other hosts via its NIC (network interface controller).

A network connects the nodes. The nodes may be arranged in some topology, which determines the network's carrying capacity and cost.

In a **Data Processing Unit (DPU)**, the NIC becomes "host-like" via the addition of processing (ypu), memory and other accelerations engines.

Node



How do we currently use SmartNICs?

Metadata management for storage systems

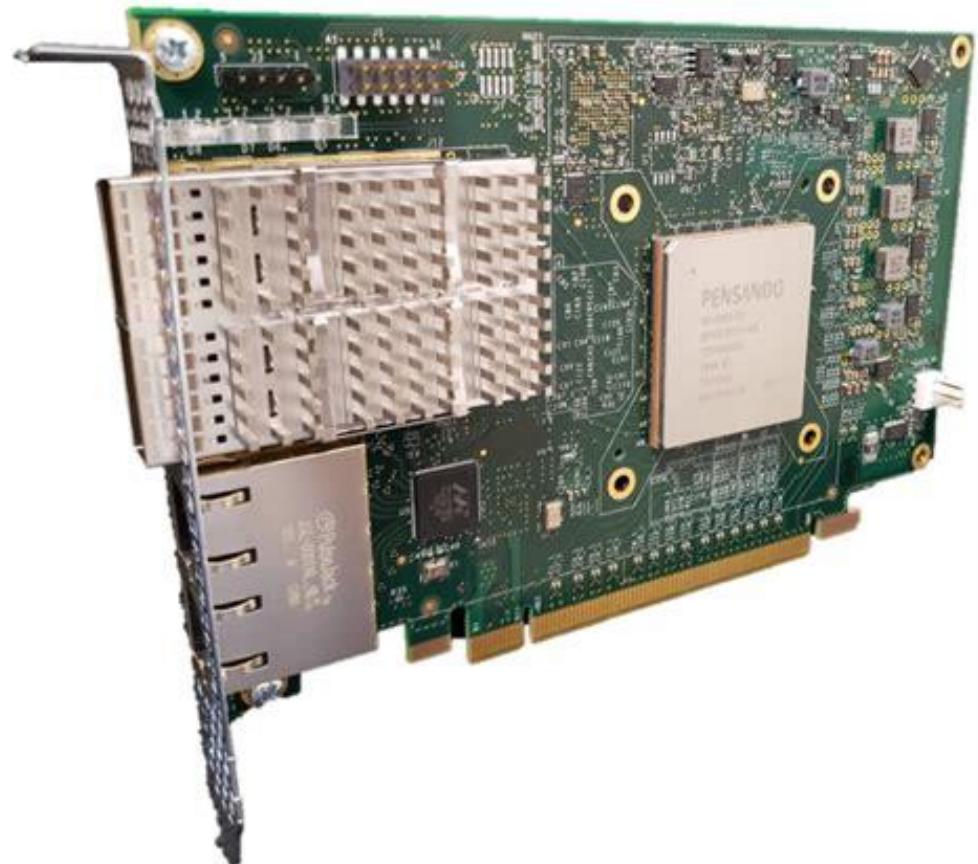
- Key part of VAST deployments

Containers for services

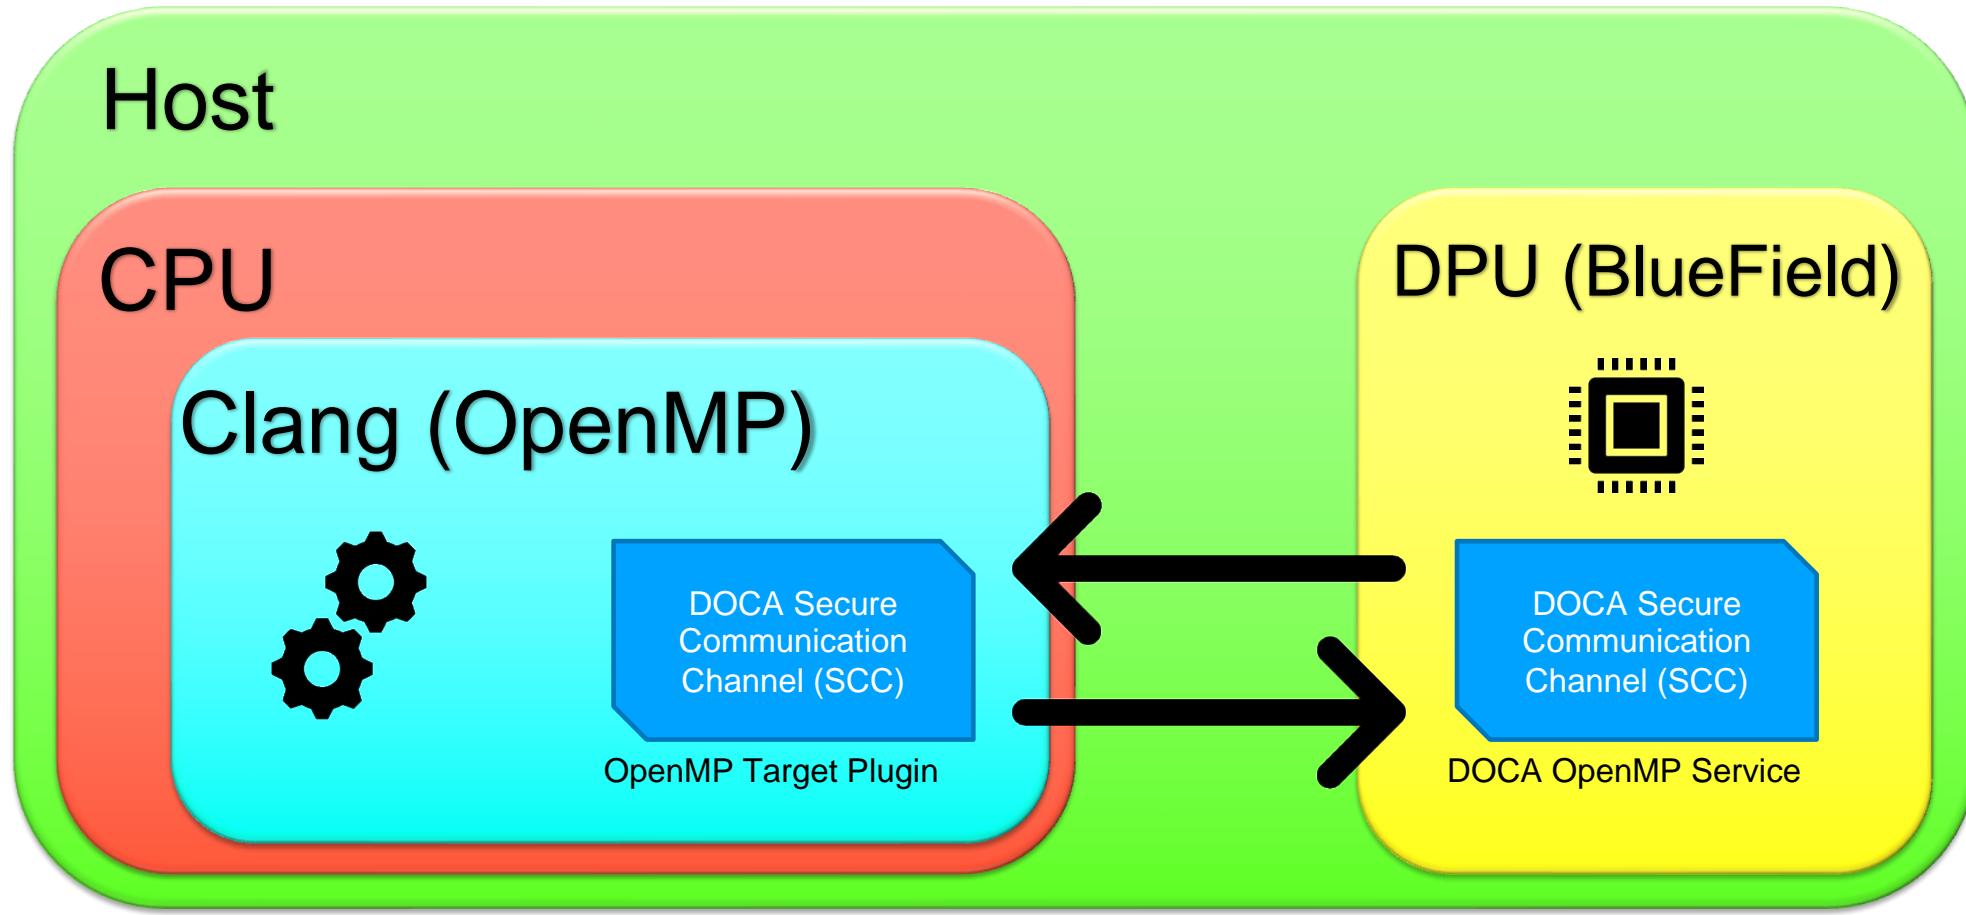
- DPU environment runs Linux and Kubernetes and can be deployed in a “Root of Trust” mode

HPC experimentation and research

- OpenMP Offloading
- Algorithmic restructuring with BlueField SmartNICs
- Development of user-defined accelerators for FPGA-based SmartNICs

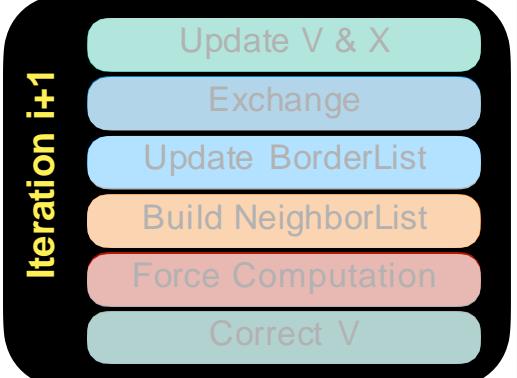
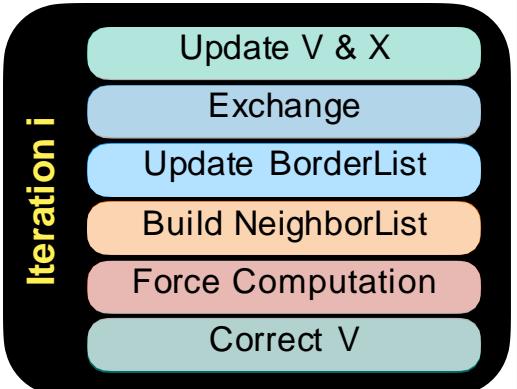
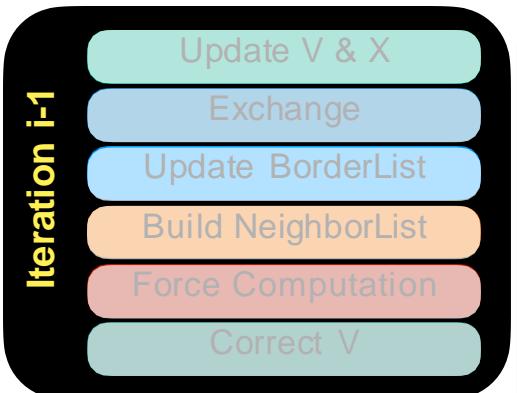


ODOS – Barcelona Supercomputing Center



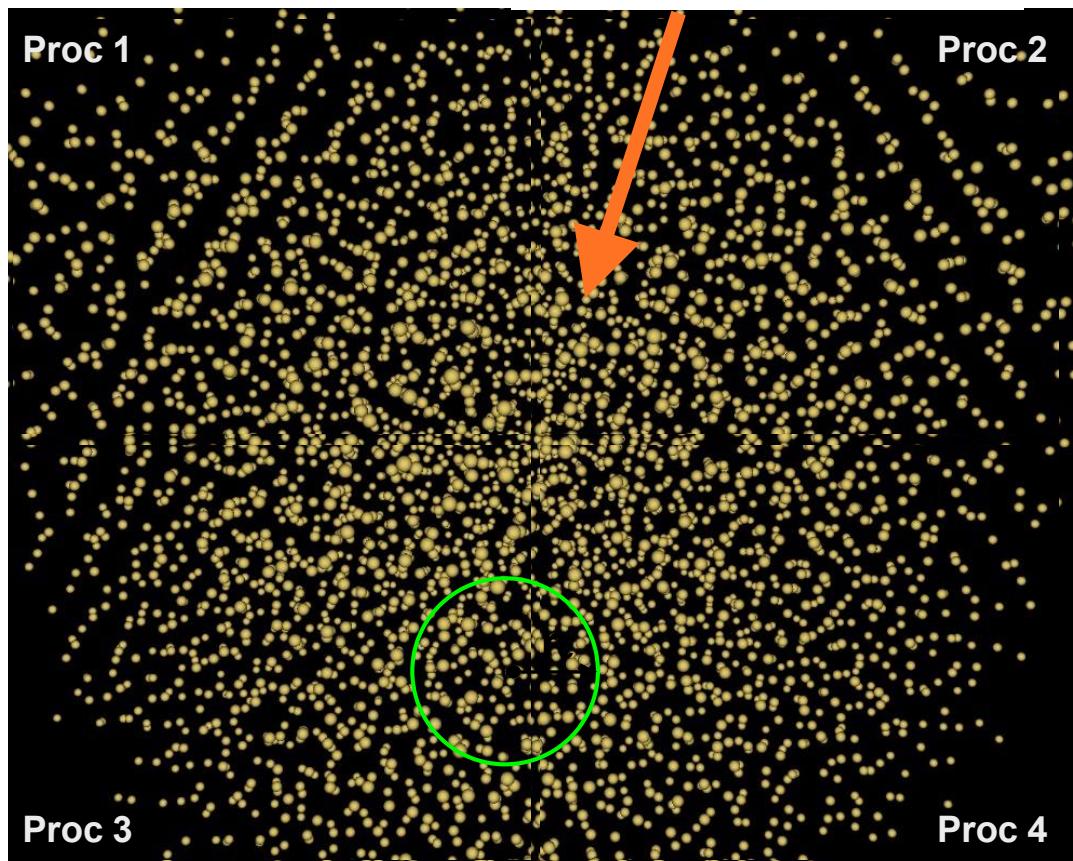
See M. Usman, S. Iserte, R. Ferrer, and A. J. Peña, "DPU Offloading Programming with the OpenMP API." LLVM-HPC23 co-located with SC23, <https://doi.org/10.1145/3624062.3624165>

Using SmartNICs with MiniMD



MiniMD is a molecular dynamics proxy-app.
It calculates the position and velocity of a set
of interacting particles in discrete time steps
(iterations).

**“Ghost zone”
(comm+overhead)**



Using SmartNICs with MiniMD

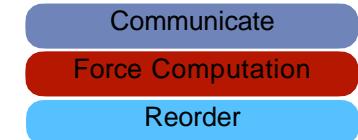
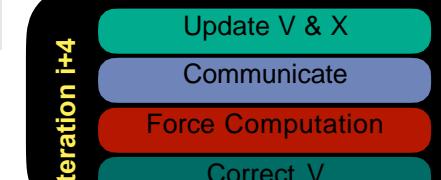
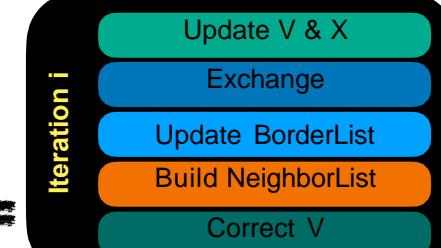
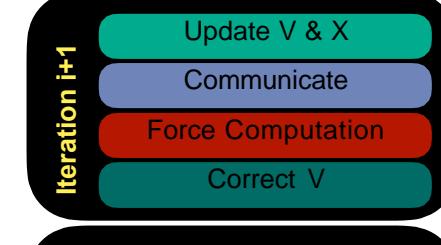
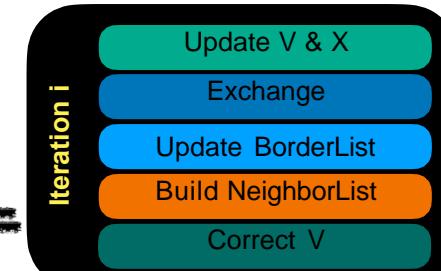
Neighbor List Generation K-1

Neighbor List Generation K

Neighbor List Generation K+1

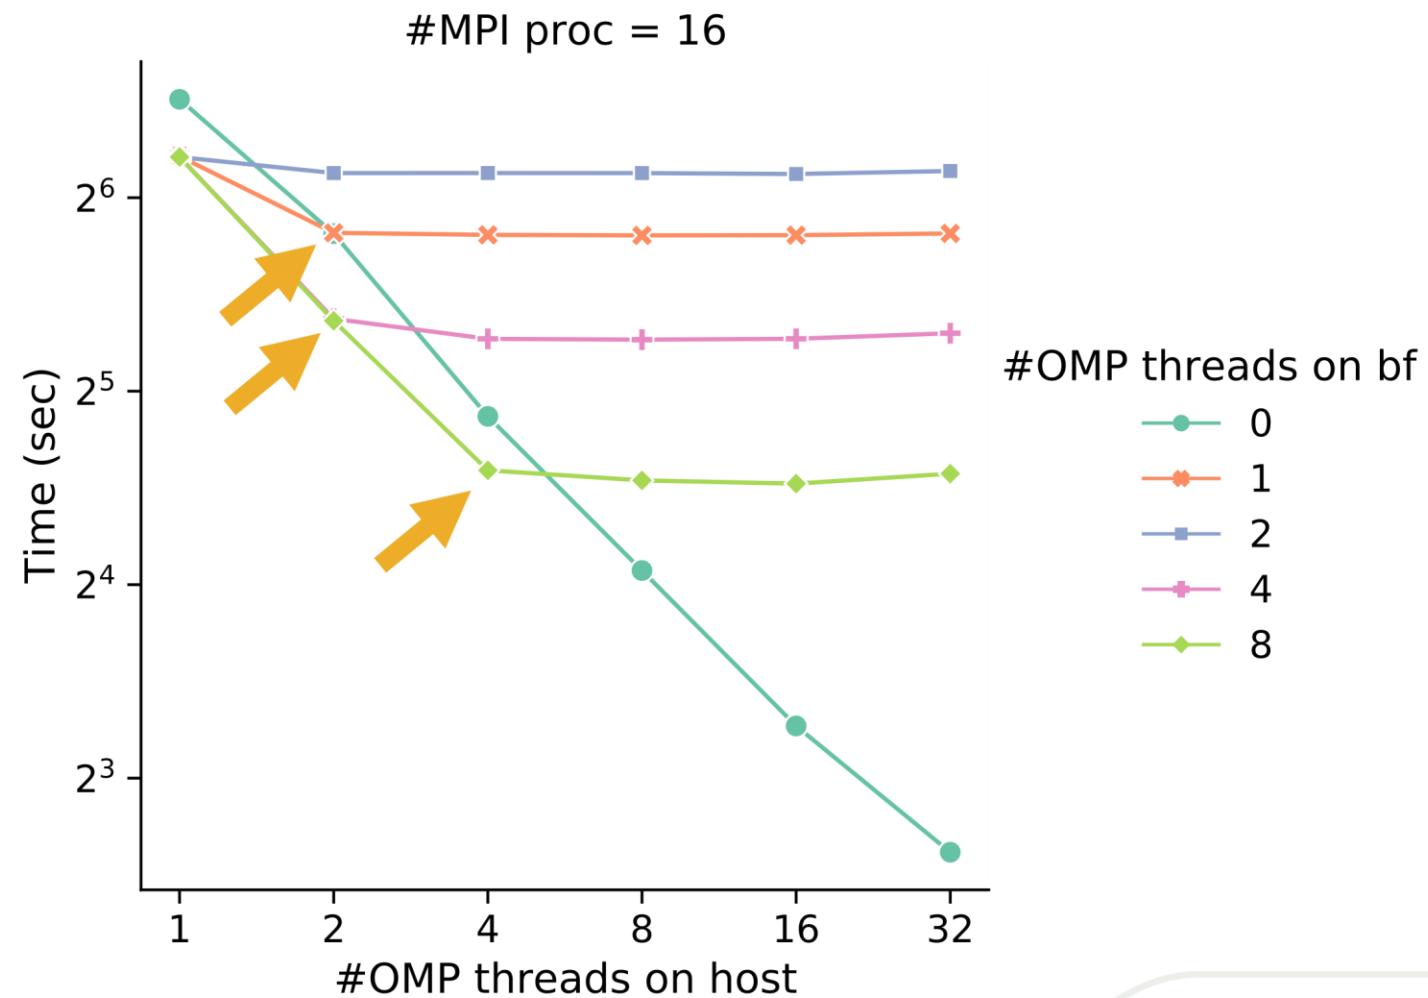
Host

DPU



MiniMD - Hybrid MPI/OpenMP performance results

- Our algorithm works best when it can completely hide the force computation time on BlueField.
- The degree of achievable overlap depends on the relative computational power of the host and BlueField.
- The **knee of each curve** indicates where the running times of neighbor-build on the host and force-compute on the BlueField are closest.



What should we be doing with DPUs?

HPC offload

- Collective communication offloading
- Offload for compute acceleration with ODOS
 - This assumes it is easy to test and show potential benefits for an application

HPC supporting services

- Containers that can help offload checkpointing or storage services for HPC applications

Secure HPC and AI computing

- Extend Trusted Execution Environments (TEEs) from one node to a larger storage/memory enclave

Prototyping Vehicle

- Test an idea for a HPC accelerator and then put it in hardware
- e.g., ZFP or compression/decompression accelerators
- Data Path Accelerators on BlueFields allow for “line-rate prototyping” without P4

What's Next?

Follow-on Effects – Evaluating the Testbed

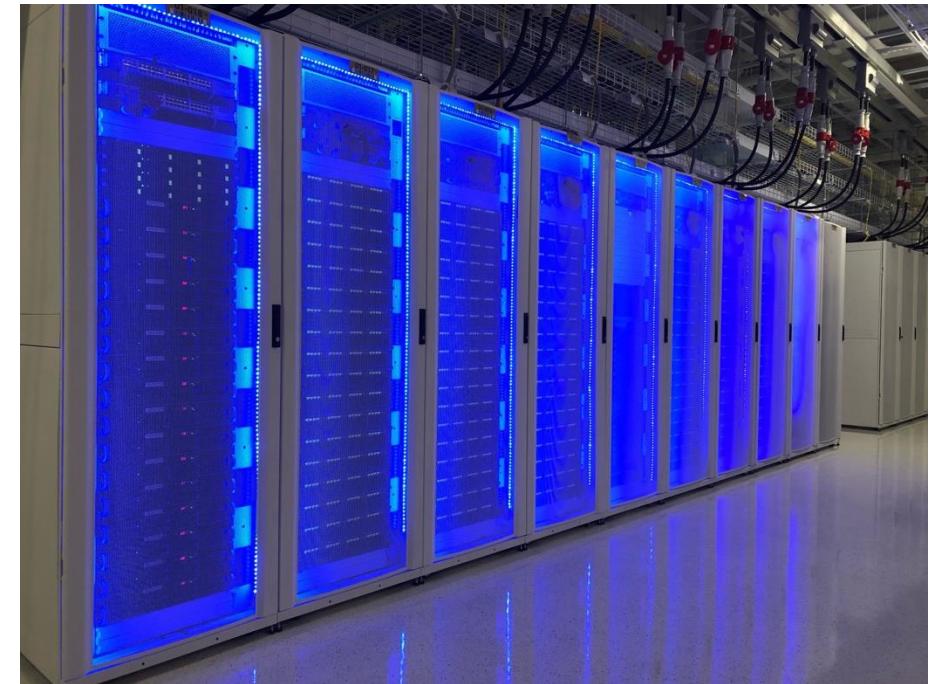
When we think about tomorrow's supercomputers, they are **highly heterogeneous**.

Testbeds like the Rogues Gallery allow us not only to evaluate far-off technologies but they are prototypes for near-term production systems

- Successful candidates migrate to our large-scale cluster, **Phoenix**, or the **Instructional Cluster Environment (ICE)**

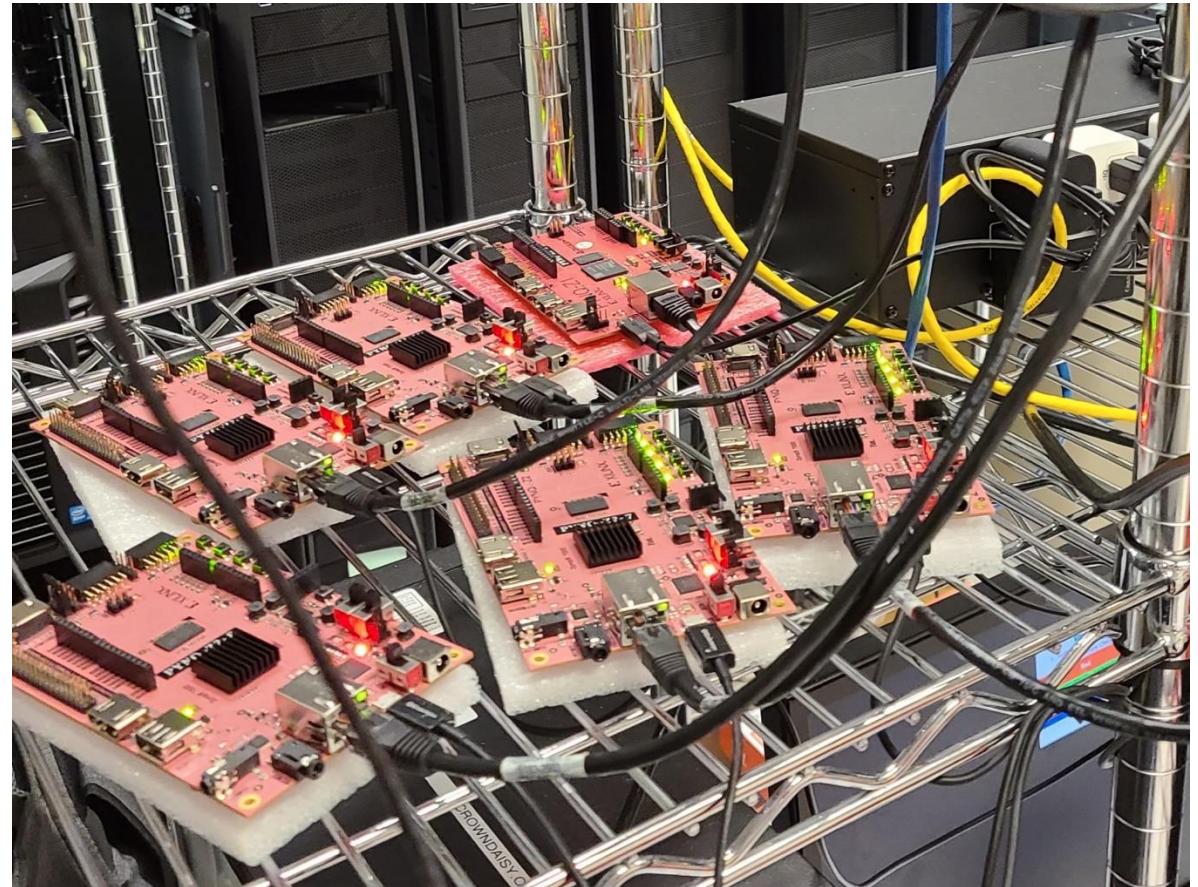
In addition to hardware, we prototype **software and tools** for deployment at scale

- Slurm features like node helpers and NSS support for launching cloud, container jobs
- Further adoption of user interfaces like Open OnDemand that make access easy so researchers can focus on the challenging systems and co-design tasks.



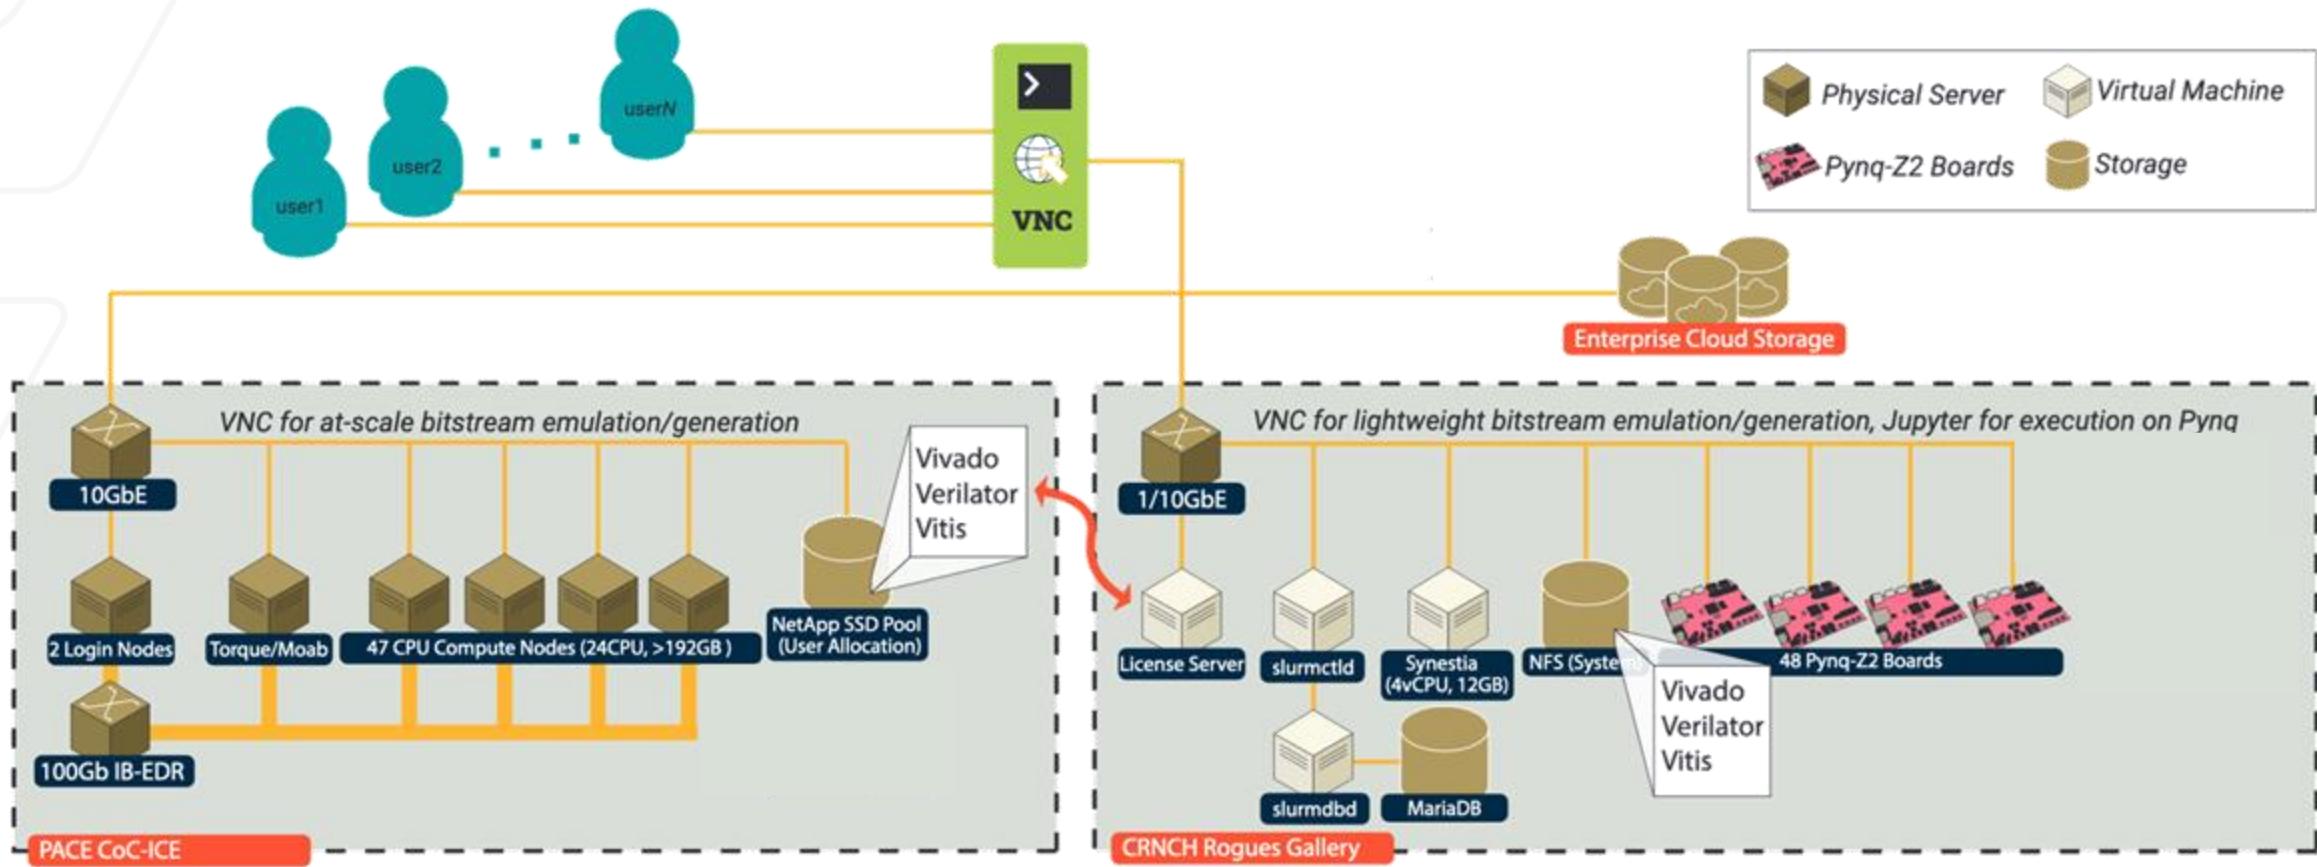
Bonus Example – Testbed + Production for Instruction

- Two courses with 90+ (CS 3220 and 40+ (ECE 8803) students
- Typical PYNQ FPGA workflows require:
 - The usage of Vivado and Vitis tools with large memory requirements
 - Hands-on interaction with a board to set up Jupyter webserver interfaces
- Supply chain issues limited the number of boards we could deploy



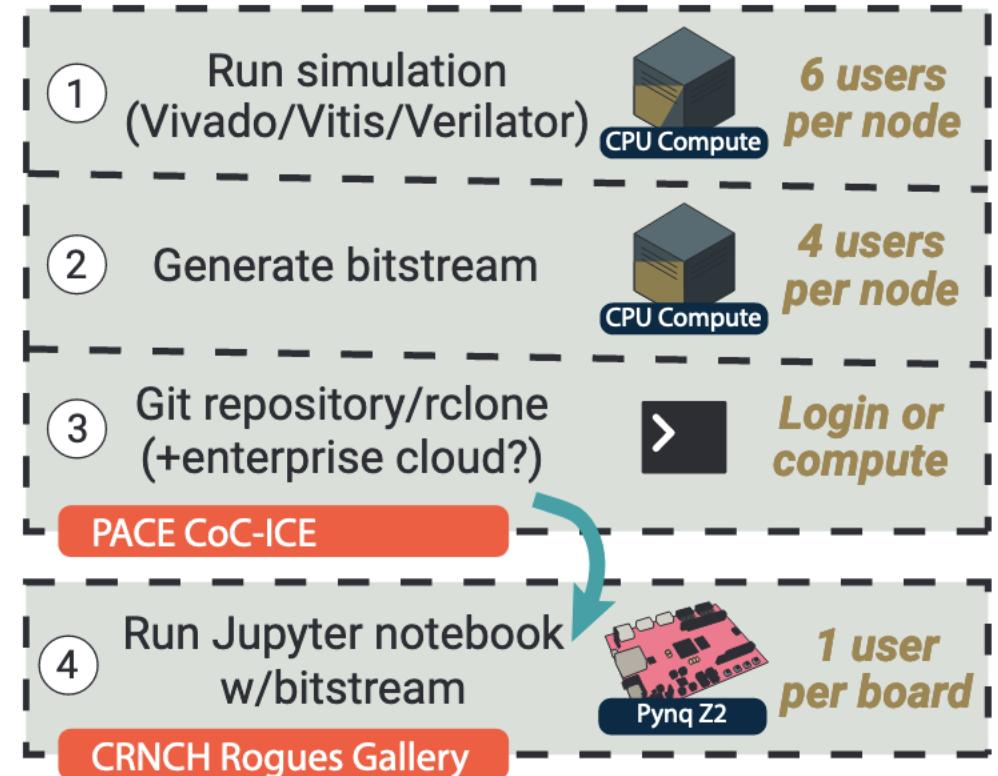
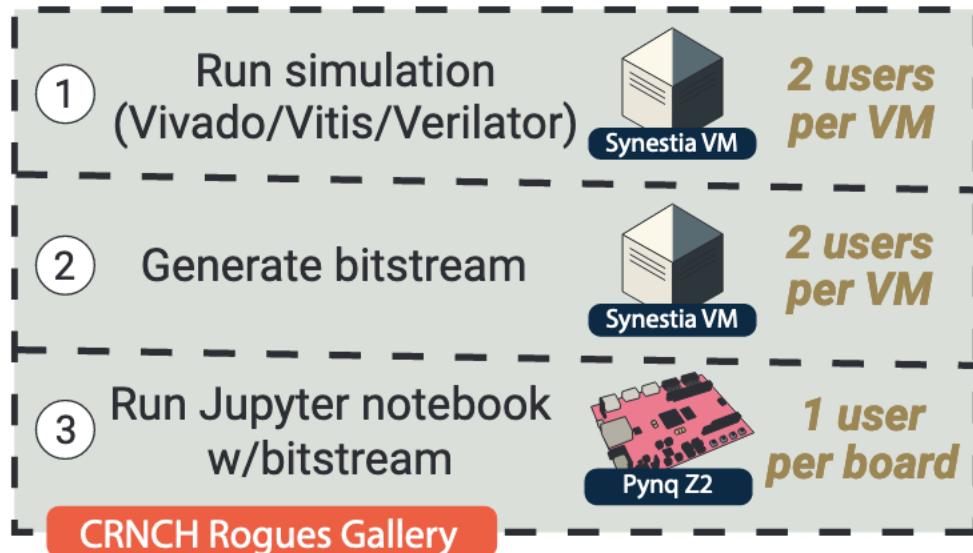
How can we best support remote FPGA development and usage at scale?

CRNCH + PACE Instructional Workflow



See our SEHET22 paper, “Enhancing HPC Education and Workflows with Novel Computing Architectures” at <https://github.com/gt-crnch-rg/rg-publications/>

Two User Workflows

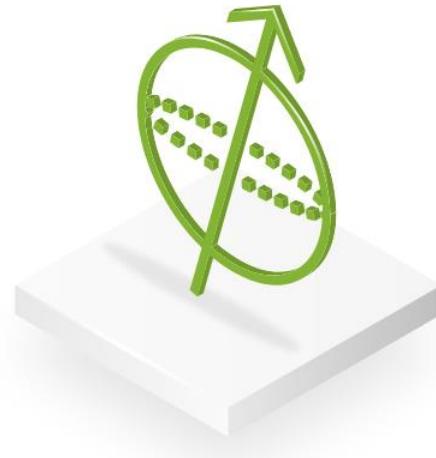


A combined workflow allows for a dramatic reduction in the bottleneck of design and initial testing

- However, there are some issues to consider related to data security
- ICE is set up to manage student data unlike research testbeds like CRNCH RG

Integration with Institutional Computing

Application portability across clusters through containerized workflows!



Chisel

PACE resources (Phoenix and ICE) provide more capacity and capability, so as problems outgrow RG, extend workflows accordingly

- Xilinx libraries on ICE for FPGA development
- Chisel/Chipyard containers on Phoenix
- CuQuantum containers and support

Heterogeneous AI Workflows

1 Common Storage

Common storage (CEDAR, IDEaS) across campus allow opportunity to migrate data/models to optimal architectures.

Toolkits for exploring alternative architectures and pipelines to evaluate appropriate hardware before production

2 Simulate and Dispatch

3 Function as a Service

Globus Compute Endpoint (funcX)/CyberShuttle as a means to target different accelerators seamlessly within a single application.

Heterogeneous User Groups

1

Temporary user access groups

- Single lab assignment, workshops, or tutorials
- Makes use of JupyterHub, Slurm job container plugins, topic-specific containers
- Requires additional security technical and policy discussions

2

Medium-term users of resources for instructional or research purposes

- Semester-long courses or research experiences
- May require longer-term access; can still use Slurm job container plugins, topic-specific containers

3

General users

- Long-term research or instructional accounts; service accounts
- Full user-level access
- Often require specialized support or limited sudo to run some commands ('nvidia-smi' or "docker run")

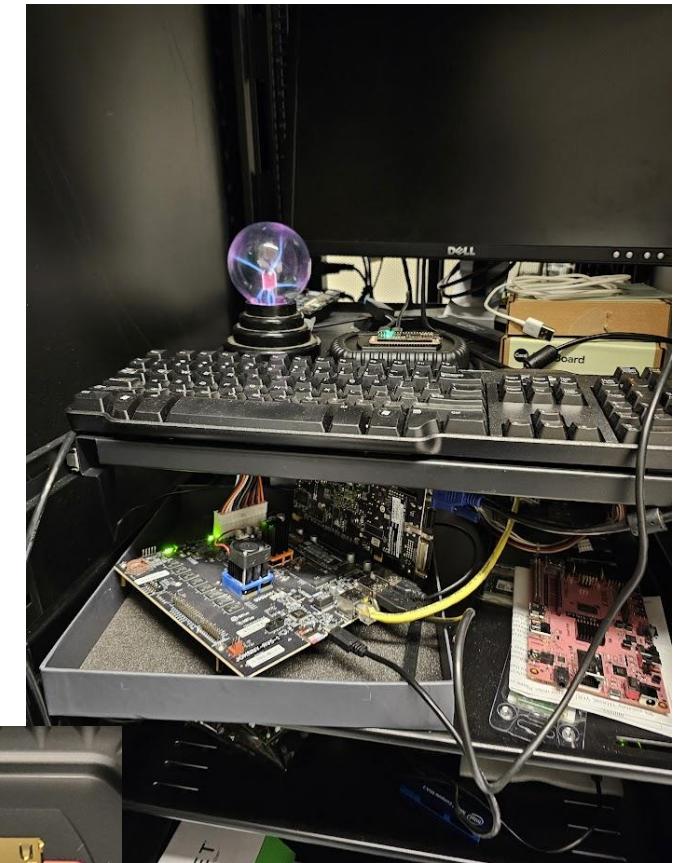
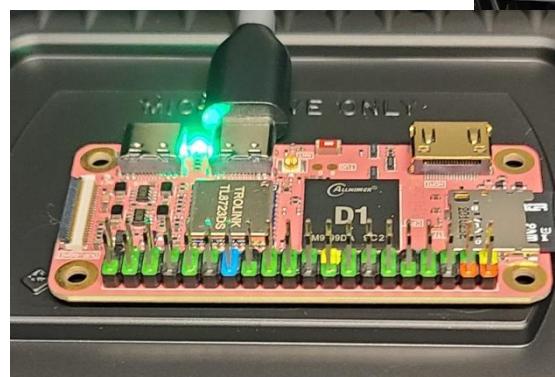
Challenges in CRNCH RG Deployment

Supporting one class of novel architecture requires an enormous investment in backend resources, training, and infrastructure

It's hard to imagine exerting the same level of effort for every new platform!

e.g., *Currently, RG uses 16 different builds of Slurm, which is only manageable with automation*

However, recognizing the patterns for typical novel architecture workflows and mapping them to common infrastructure can provide a solid basis for research

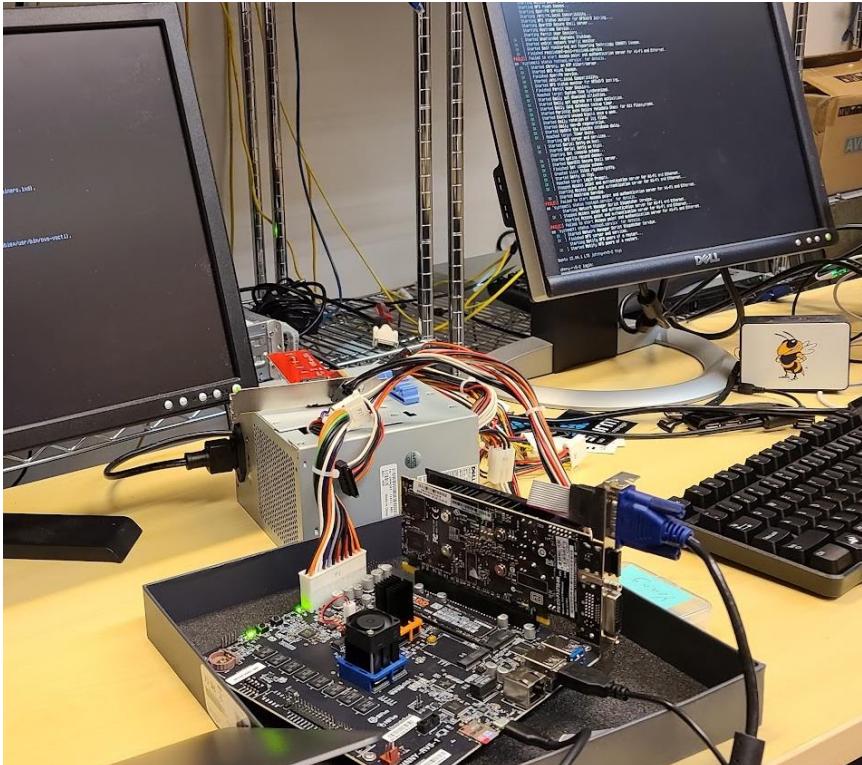
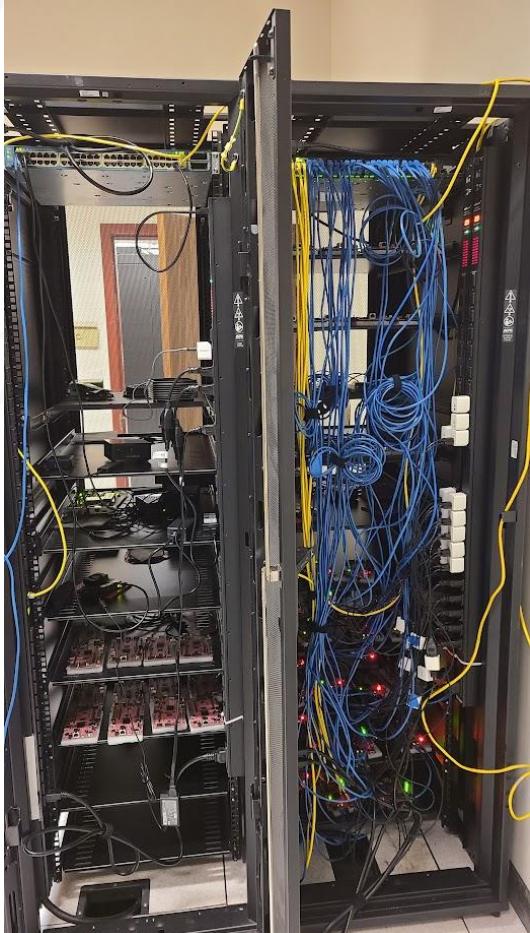


Summary

Novel architectures will be critical to understanding data movement trends and designing solutions for post-exascale systems

- Device heterogeneity is here to stay and it may be critical for future performance!
- Testbeds can be a key vehicle for both this research and related instruction and training
- Getting the proper support for a platform is critical to its success
 - Each “candidate architecture” needs enthusiast users and dedicated researchers invested in research success
 - Courses and applications that can make use of a unique feature are also good targets for development and usage of novel architectures
 - Training and high-level tools are vital to reduce barriers to usage

Questions?



Project site: crnch-rg.cc.gatech.edu
GH Presence: <https://github.com/gt-crnch-rg>

Apply for access: <https://crnch-rg.cc.gatech.edu/request-rogues-gallery-access/>

Many Thanks to the CRNCH and RG Contributors

NSF Project Co-PIs



Tom Conte



Jennifer Hasler

Senior Personnel



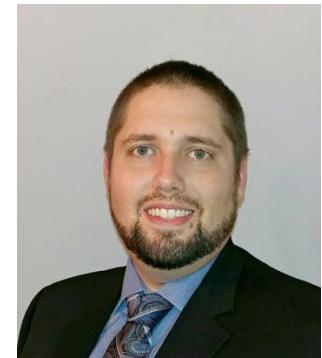
Alex Daglis



Sterling Peet



Will Powell



Aaron Jezghani (PACE)



Ada Gavrilovska



Rich Vuduc

Acknowledgements

This work is supported by the NSF Rogues Gallery testbed grant (CNS-2016701)

This research used resources provided by the Darwin testbed at Los Alamos National Laboratory (LANL) which is funded by the Computational Systems and Software Environments subprogram of LANL's Advanced Simulation and Computing program (NNSA/DOE).

Kevin Sheridan, Galen Shipman, and Jered Dominguez-Trujillo acknowledge support by the National Nuclear Security Administration. Los Alamos National Laboratory is operated by Triad National Security, LLC for the U.S. Department of Energy under contract 89233218CNA000001; LA-UR-24-24856.